

CASSANDRA
SUMMIT 2014

| EUROPE



London, U.K.
December 3 - 4

#CassandraSummit
Park Plaza Riverbank Hotel



Training Day | *December 3rd*

Beginner Track

- *Introduction to Cassandra*
- *Introduction to Spark, Shark, Scala and Cassandra*

Advanced Track

- *Data Modeling*
- *Performance Tuning*

Conference Day | *December 4th*

Cassandra Summit Europe 2014 will be the single largest gathering of Cassandra users in Europe. Learn how the world's most successful companies are transforming their businesses and growing faster than ever using Apache Cassandra.

<http://bit.ly/cassandrasummit2014>

EUROPE'S LARGEST GATHERING OF CASSANDRA DEVELOPERS

DECEMBER 3 - 4, 2014 | LONDON, U.K. | THE PARK PLAZA RIVERBANK HOTEL

Twitter: [@CassandraEurope](#) | [#CassandraSummit](#)



Introduction to Cassandra for Java Developers

Why, What and How.

Johnny Miller, Solutions Architect

@CyanMiller



www.linkedin.com/in/johnnymiller





Training Day

December 3rd

Beginner Track

- *Introduction to Cassandra*
- *Introduction to Spark, Shark, Scala and Cassandra*

Advanced Track

- *Data Modeling*
- *Performance Tuning*

Conference Day

December 4th

Cassandra Summit Europe 2014 will be the single largest gathering of Cassandra users in Europe. Learn how the world's most successful companies are transforming their businesses and growing faster than ever using Apache Cassandra.

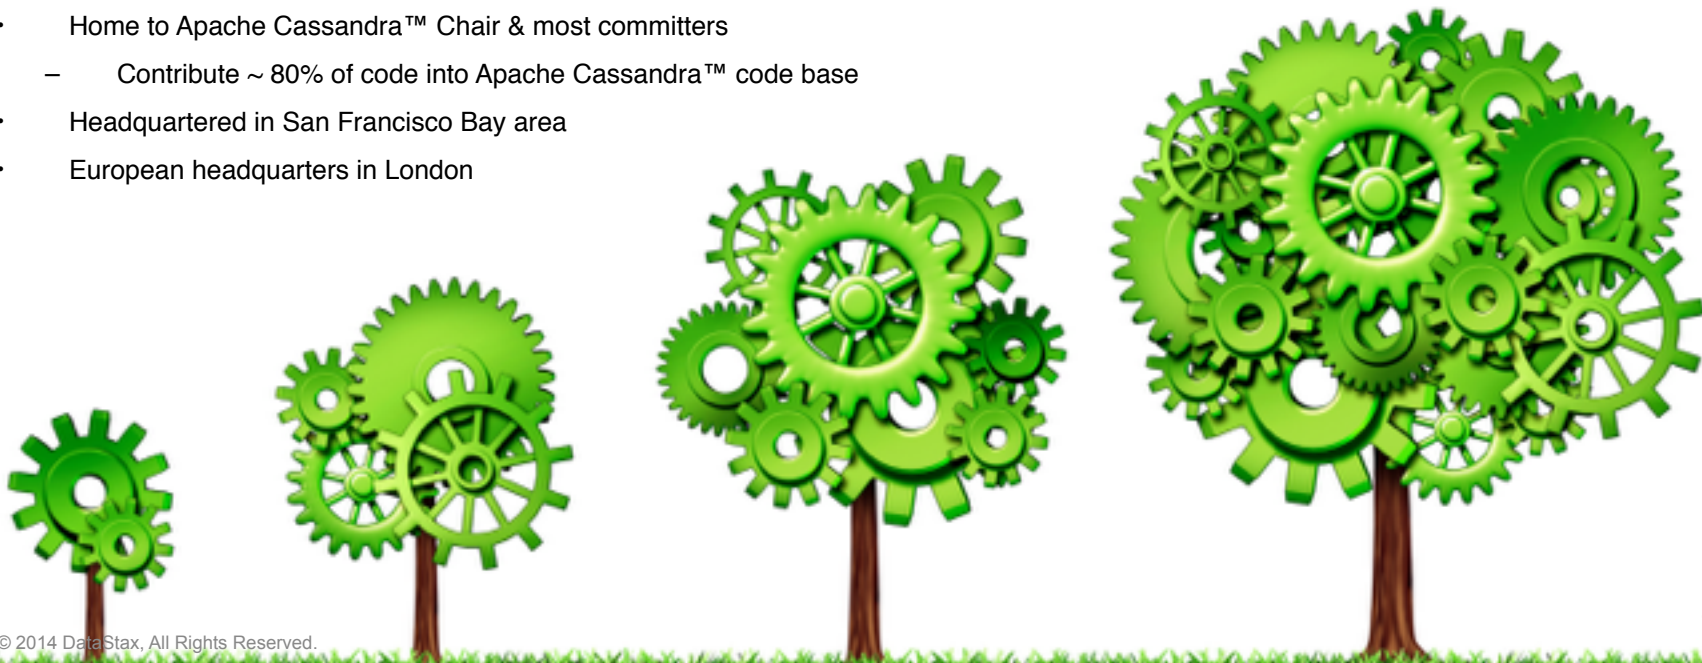
<http://bit.ly/cassandrasummit2014>

EUROPE'S LARGEST GATHERING OF CASSANDRA DEVELOPERS

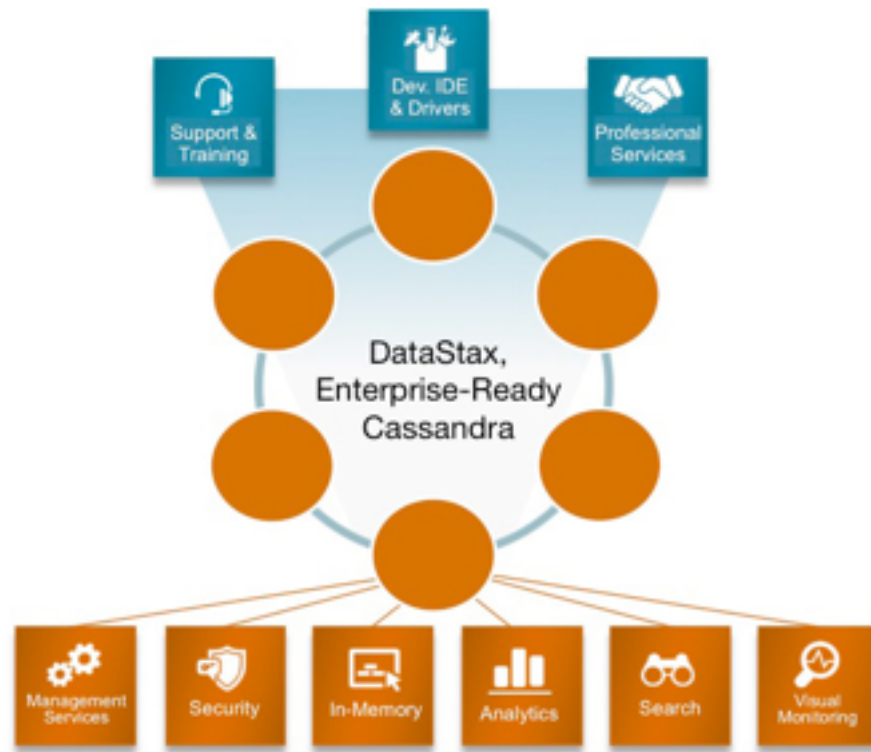
DECEMBER 3 - 4, 2014 | LONDON, U.K. | THE PARK PLAZA RIVERBANK HOTEL

Twitter: [@CassandraEurope](#) | [#CassandraSummit](#)

- Founded in April 2010
- We drive Apache Cassandra™
- 300+ employees
- Home to Apache Cassandra™ Chair & most committers
 - Contribute ~ 80% of code into Apache Cassandra™ code base
- Headquartered in San Francisco Bay area
- European headquarters in London



DataStax Enterprise



www.datastax.com

DataStax Enterprise is free for startups!!!

DataStax  Startups

- **Unlimited, free use of the software in DataStax Enterprise.**
- No limit on number of nodes or other hidden restrictions.
- If you're a startup, it's free!

www.datastax.com/startups

Training

- Checkout the DataStax Academy for **free** online training
 - <http://academy.datastax.com>
- Public courses
- On-site training



www.datastax.com/training

We are hiring!

www.datastax.com/careers
@DataStaxCareers

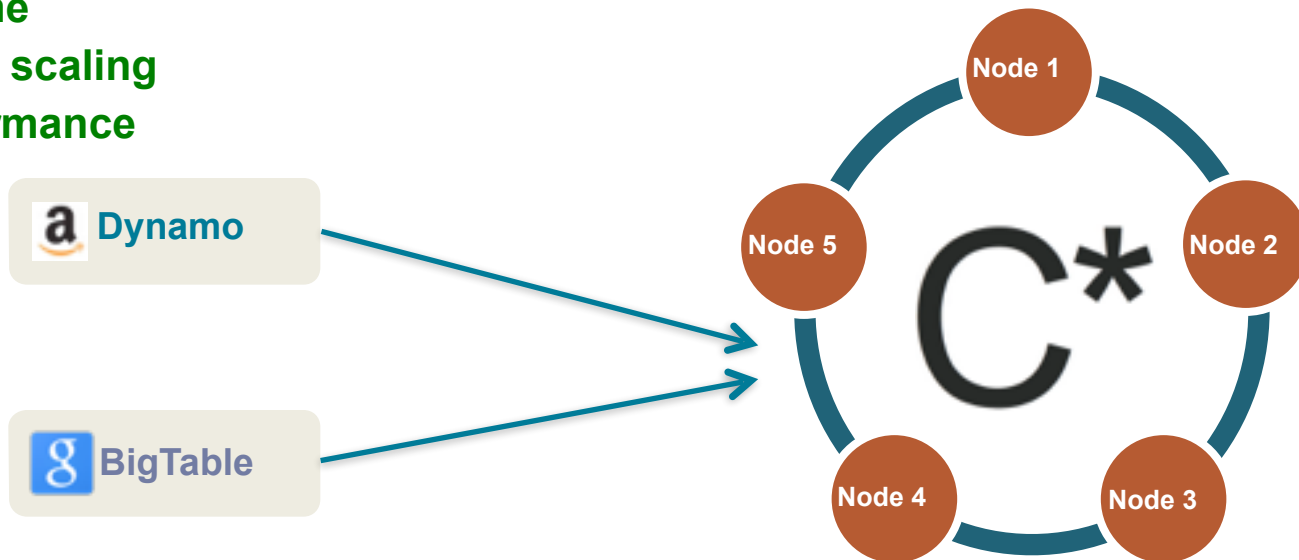


Why Cassandra?



Apache Cassandra

- Cassandra is a massively scaleable **open source** NoSQL **database**
- **100% uptime**
- **Predictable scaling**
- **High Performance**



BigTable:

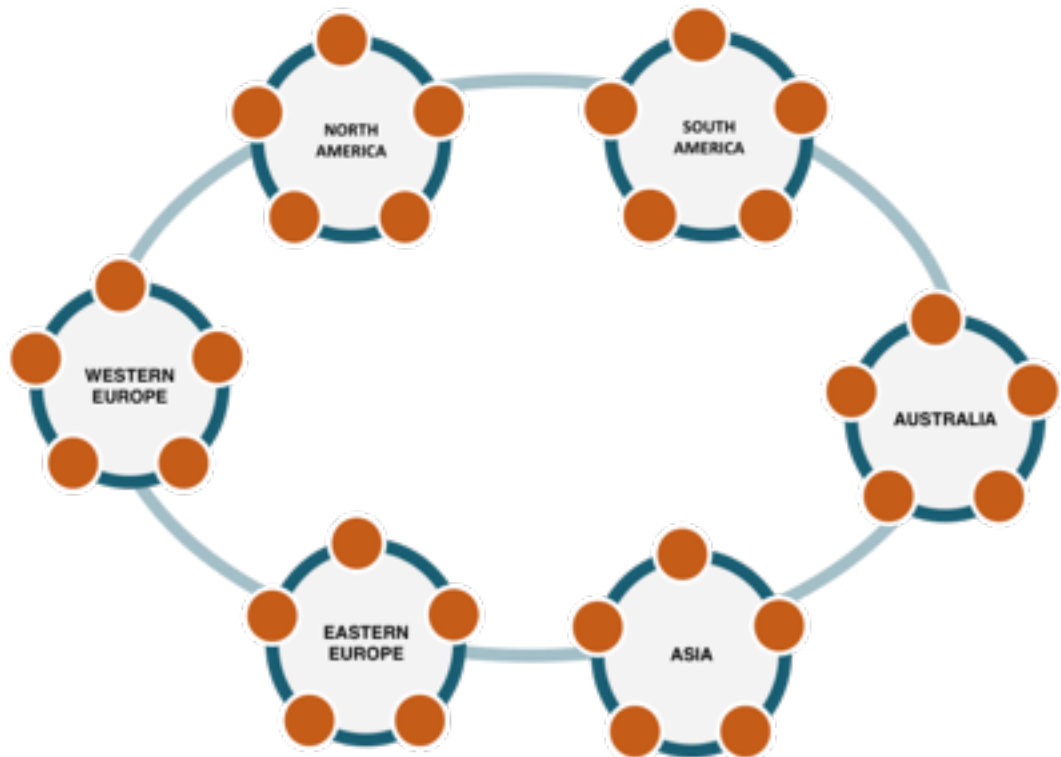
<http://research.google.com/archive/bigtable-osdi06.pdf>

Dynamo:

<http://www.allthingsdistributed.com/files/amazon-dynamo-sosp2007.pdf>

Cassandra Core Values

- **Ease of use**
- **Massive scalability**
- **High performance**
- **Always Available**



Availability and Speed Matters for online apps!

- UK retailers **lost 8.5 billion last year to slow web sites**, which is 1 million for every 10 million in online sales
- Over half of all web users expect a **response time of 2 seconds or less**
- A 1 second delay causes a nearly 10% reduction in customer interactions
- A 1 second decrease in Amazon page load time costs the company \$1.6 billion in sales

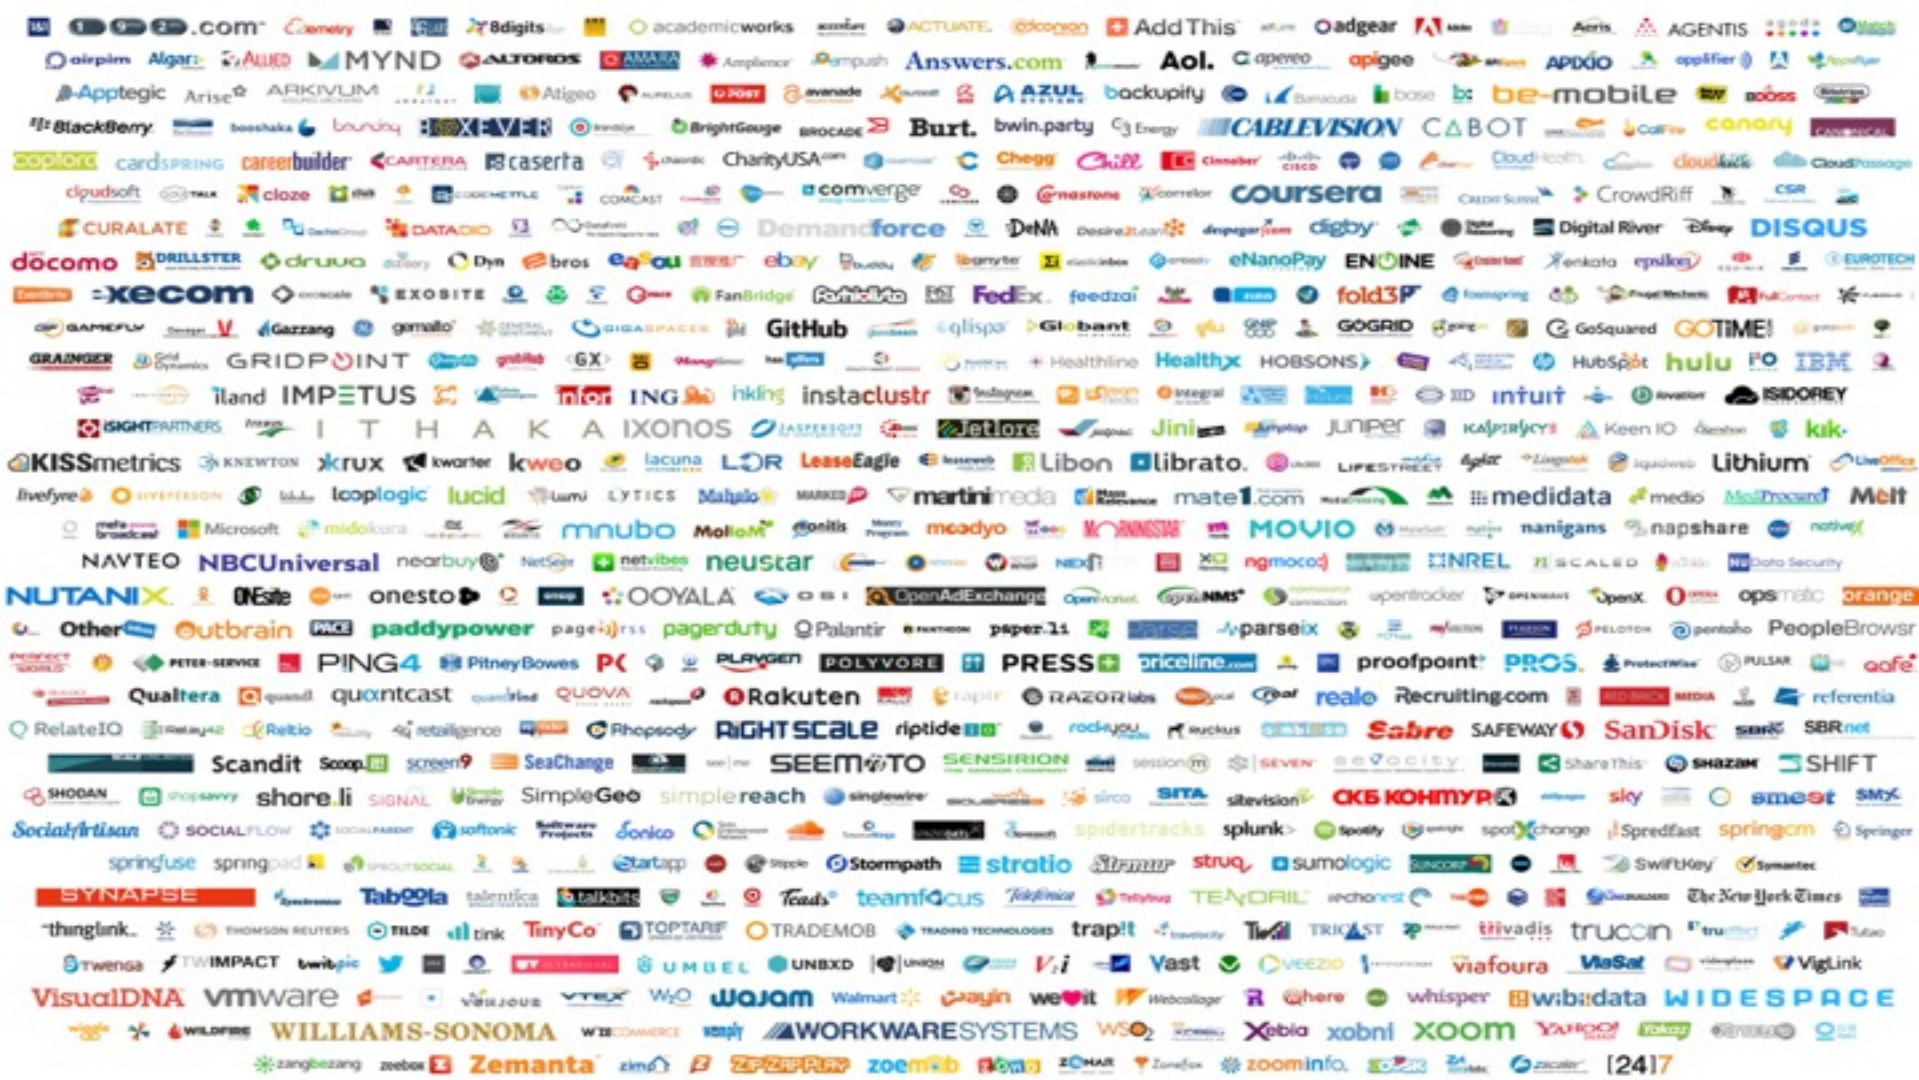


Adoption

Rank	Last Month	DBMS	Database Model	Score	Changes
1.	1.	Oracle	Relational DBMS	1466.91	-3.95
2.	2.	MySQL	Relational DBMS	1297.14	+15.92
3.	3.	Microsoft SQL Server	Relational DBMS	1208.87	-33.62
4.	4.	PostgreSQL	Relational DBMS	255.79	+5.94
5.	5.	MongoDB	Document store	240.98	+3.63
6.	6.	DB2	Relational DBMS	197.03	-9.39
7.	7.	Microsoft Access	Relational DBMS	140.48	+0.86
8.	8.	SQLite	Relational DBMS	92.61	+3.74
9.	↑	10. Cassandra	Wide column store	87.86	+5.96
10.	↓	9. Sybase ASE	Relational DBMS	85.42	-0.75
11.	↑	12. Solr	Search engine	75.77	+6.74
12.	↓	11. Redis	Key-value store	74.60	+3.80
13.		13. Teradata	Relational DBMS	66.15	+0.77
14.		14. FileMaker	Relational DBMS	52.63	+0.57
15.		15. HBase	Wide column store	45.02	+3.11
16.		16. Elasticsearch	Search engine	41.52	+2.72
17.		17. Informix	Relational DBMS	33.45	+0.29
18.		18. Memcached	Key-value store	31.59	+0.60
19.		19. Hive	Relational DBMS	31.40	+0.59
20.		20. Splunk	Search engine	29.76	+2.15
21.		21. CouchDB	Document store	25.97	+1.84
22.		22. Neo4j	Graph DBMS	24.22	+1.31
23.		23. SAP HANA	Relational DBMS	23.18	+0.86
24.		24. Couchbase	Document store	19.67	+2.13



<http://db-engines.com/en/ranking>
September 2014



Cassandra works for small to huge deployments.

- Cassandra footprint @ Netflix
- 80+ Clusters
- 2500+ nodes
- 4 Data Centres (Amazon Regions)
- > 1 Trillion transactions per day



See: <http://www.datastax.com/resources/casestudies/netflix>

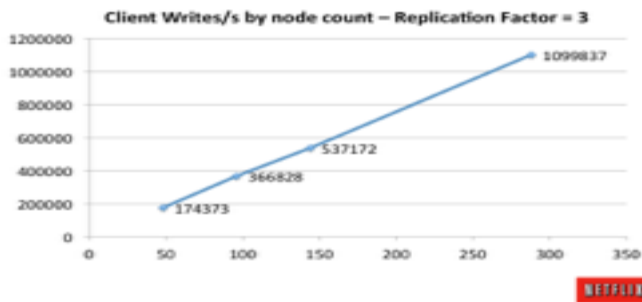
Performance & Scale

“In terms of scalability, there is a clear winner throughout our experiments. Cassandra achieves the highest throughput for the maximum number of nodes in all experiments with a linear increasing throughput.”

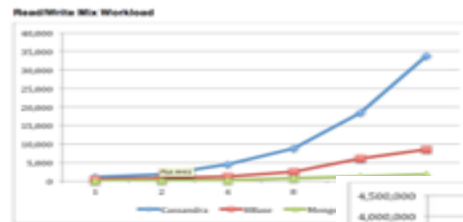
Solving Big Data Challenges for Enterprise Application Performance Management, Tilman Rable, et al., August 2012. Benchmark paper presented at the Very Large Database Conference, 2012.
http://vldb.org/pvldb/vol5/p1724_tilmanrabl_vldb2012.pdf



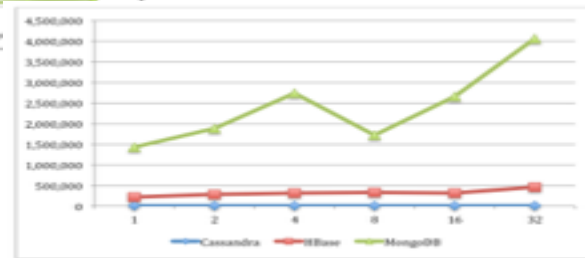
Netflix Cloud Benchmark... Scale-Up Linearity



End Point Independent NoSQL Benchmark Highest in throughput...



Lowest in latency...



Availability



Nathan Milford

@NathanMilford



Follow

Man, I just love Cassandra. Lost a data center Hurricane Sandy, nodes came up and started working with no pain.



Jake Luciani

@tjake



Following

We've lost multiple SSDs in our **#Cassandra** cluster and the JBOD support in C* 1.2 kept the nodes running while we swapped them out live!



Aaron Turner

@synfinatic



DATASTAX
comSysto
DATA AND STORAGE

took me 10hrs to notice a **#cassandra** node had a hw failure because everything just kept working. **#sweet**



Eric Florenzano

@ericflo

Following



"Cassandra ... dealt with the loss of one third of its regional nodes without any loss of data or availability."
techblog.netflix.com/2012/07/lesson... - Nice!



Bill de hÓra

@dehora



Coming to the conclusion that **#cassandra** is kind of indestructible. "Robust" doesn't do it justice.

Ideal Cassandra Use Cases

Product catalogs and playlists

- Collection of items

Internet of things

- Sensor data

Messaging

- Emails, instant messages, alerts, comments

Recommendation and personalization

- Trend information

Fraud detection

- Data patterns

Time-series and time-ordered data

- Data with time dimension



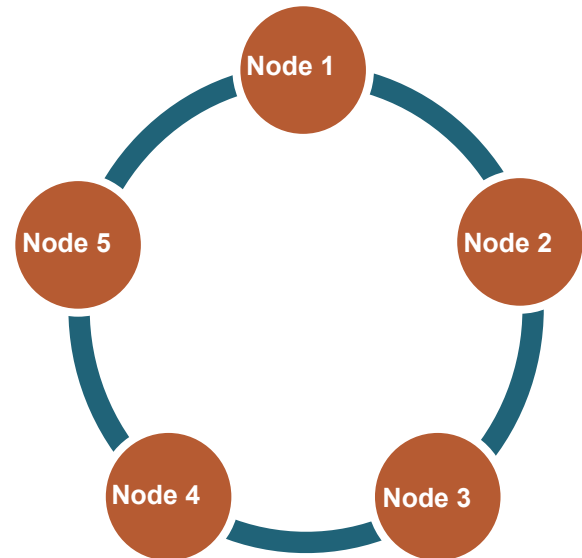
<http://planetcassandra.org>

What is Apache Cassandra?



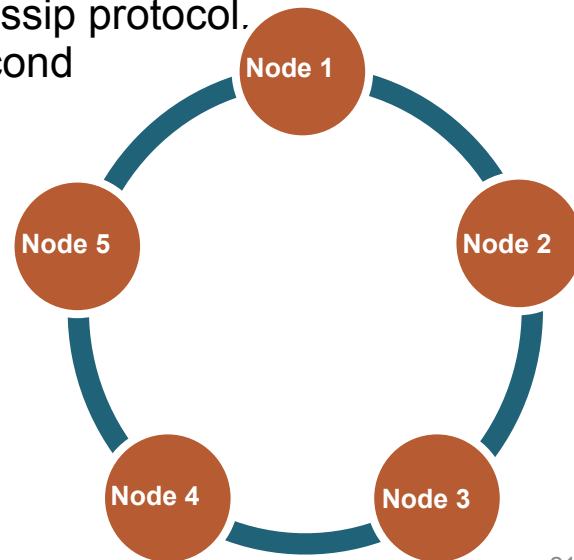
Overview

- Cassandra was designed with the understanding that system/hardware failures can and do occur
- Peer-to-peer, distributed system
- **All nodes the same**
- **Data partitioned among all nodes** in the cluster
- Custom data replication to ensure fault tolerance
- Read/Write-anywhere and across data centres



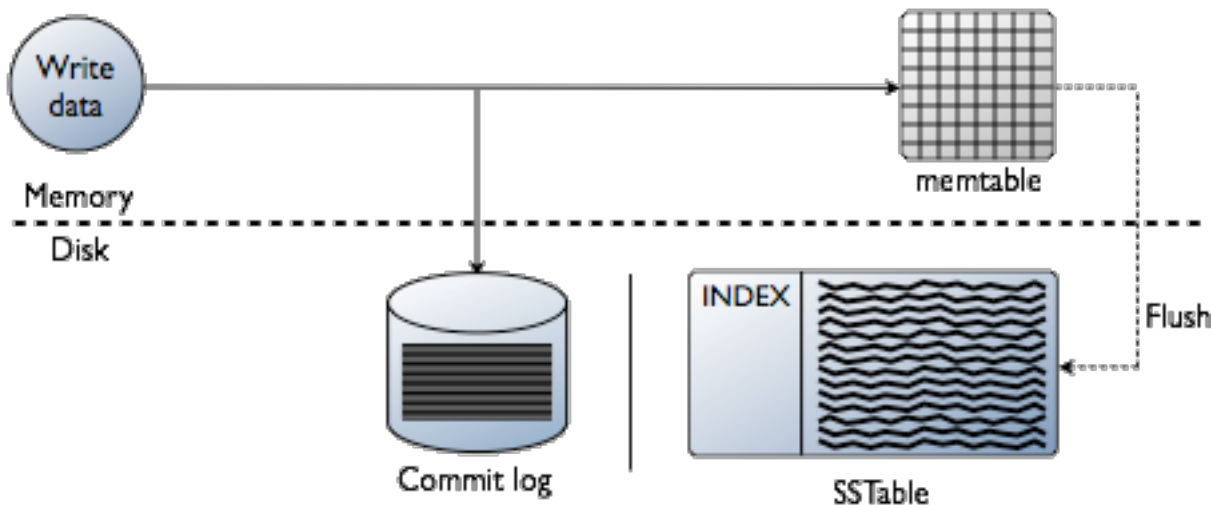
Cassandra > 1 Server

- All nodes participate in a cluster
- Add or remove as needed
- All nodes the same – **masterless** with no single point of failure
- Each node communicates with each other through the Gossip protocol. which exchanges information across the cluster every second
- **Data partitioned among all nodes in the cluster**
- **More capacity? Add a server!**
- **More throughput? Add a server!**

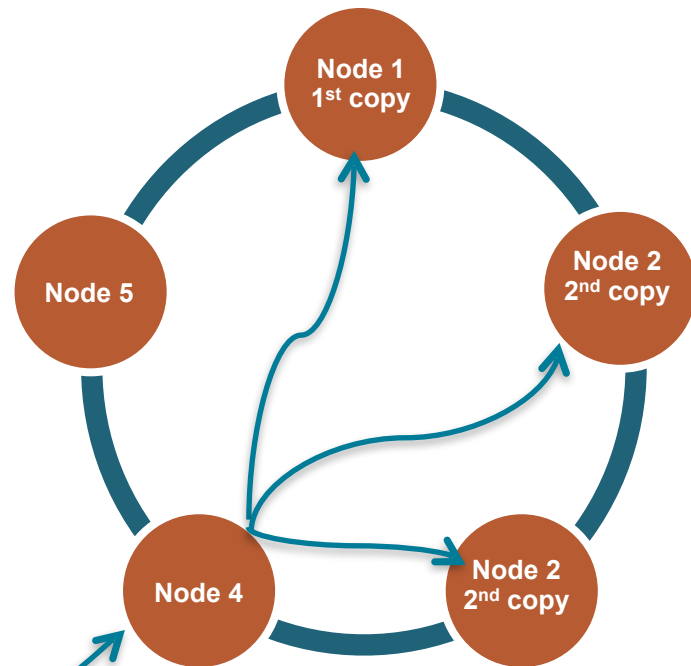


Write Path on individual server

- **A commit log** is used on each node to capture write activity. Data durability is assured
- **Data also written to an in-memory structure** (memtable) and then to disk once the memory structure is full (an SStable)



- Client reads or writes **to any node**
- Node **coordinates** with others
- Data read or replicated in **parallel**
- Replication factor (RF): How many copies of your data?
- RF = 3 in this example
- Each node is storing 60% of the clusters total data i.e. 3/5

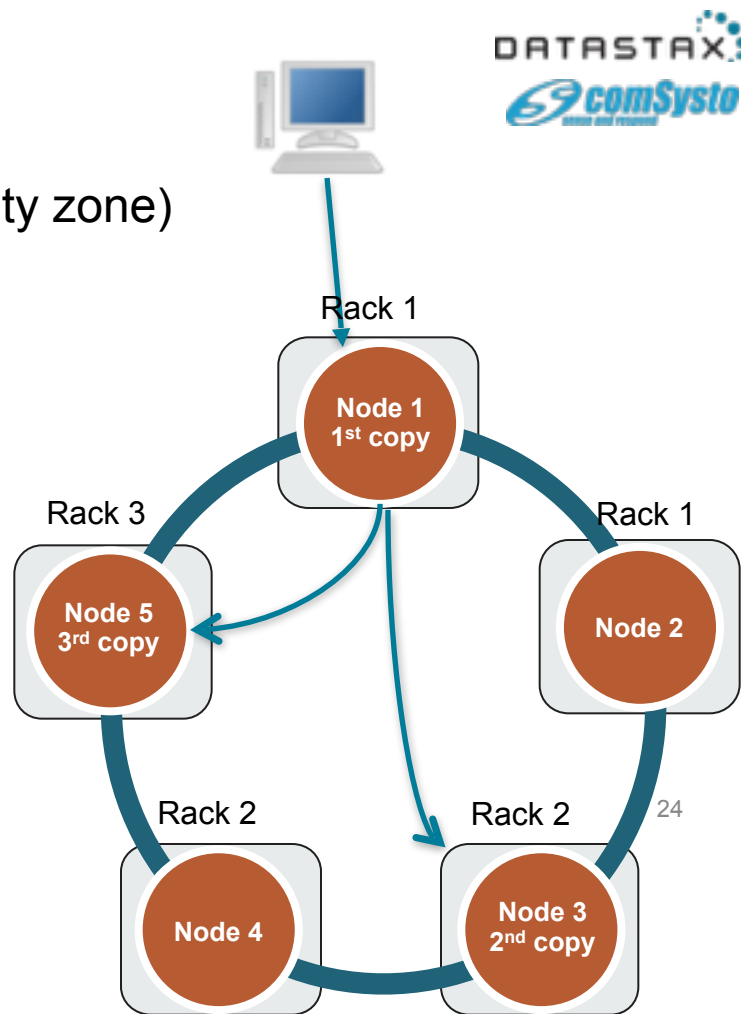


Handy Calculator: <http://www.ecyrd.com/cassandrascalculator/>



Rack Aware

- Cassandra is aware of which rack (or availability zone) each node resides in.
- **It will attempt to place each data copy in a different rack.**
- RF = 3 in this example



Data Distribution and Replication

- Cassandra is **designed as a peer-to-peer system that makes copies of the data** and distributes the copies among a group of nodes
- **Data is organized by table and identified by a primary key.** The primary key determines which node the data is stored on.
- **Each node is assigned data based on the hash value of the primary key** and the range of hash values associated with that node.



- **A partitioner determines how data is distributed across the nodes in the cluster.**
- Cassandra has three:
 - Murmur3Partitioner (default)
 - RandomPartitioner
 - ByteOrderedPartitioner
(don't use this unless you know what your doing and have a very good reason to do this)



Murmer3 Example

Primary Key

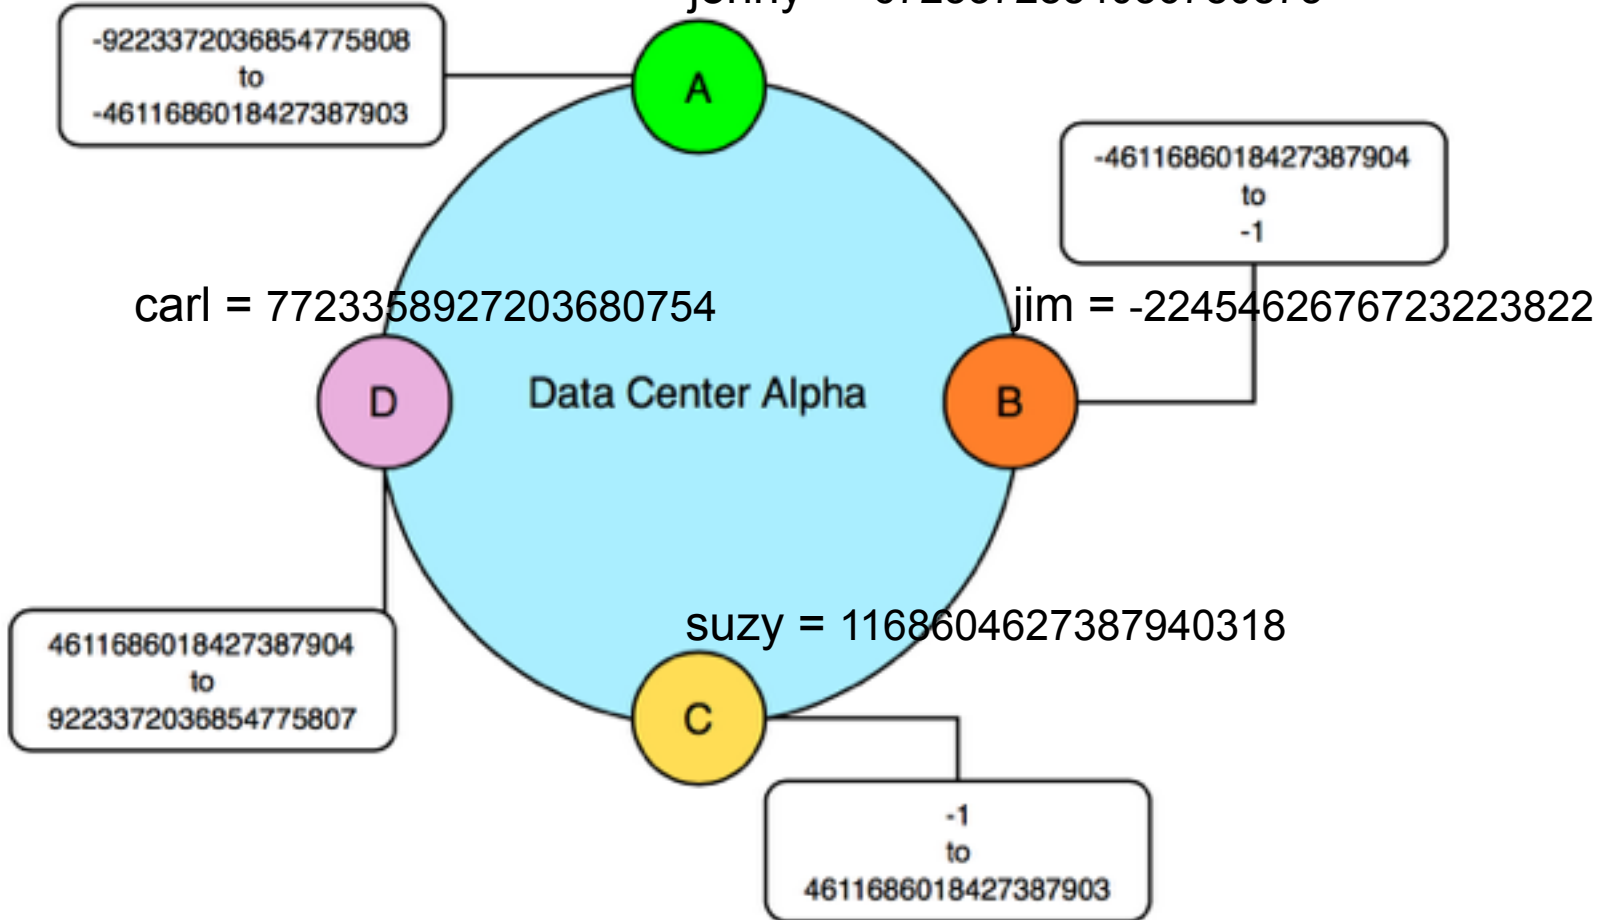
- Data:**

jim	age: 36	car: ford	gender: M
carol	age: 37	car: bmw	gender: F
johnny	age: 12	gender: M	
suzy:	age: 10	gender: F	

- Murmer3 Hash Values:**

Primary Key	Murmur3 hash value
jim	-2245462676723223822
carol	7723358927203680754
johnny	-6723372854036780875
suzy	1168604627387940318

johnny = -6723372854036780875



Murmer3 Example

Data is distributed as:

Node	Start range	End range	Primary key	Hash value
A	-922337203685477580 8	-4611686018427387903	johnny	-6723372854036780 875
B	-461168601842738790 4	-1	jim	-2245462676723223 822
C	0	4611686018427387903	suzy	1168604627387940 318
D	4611686018427387904	9223372036854775807	carol	7723358927203680 754

Replicas

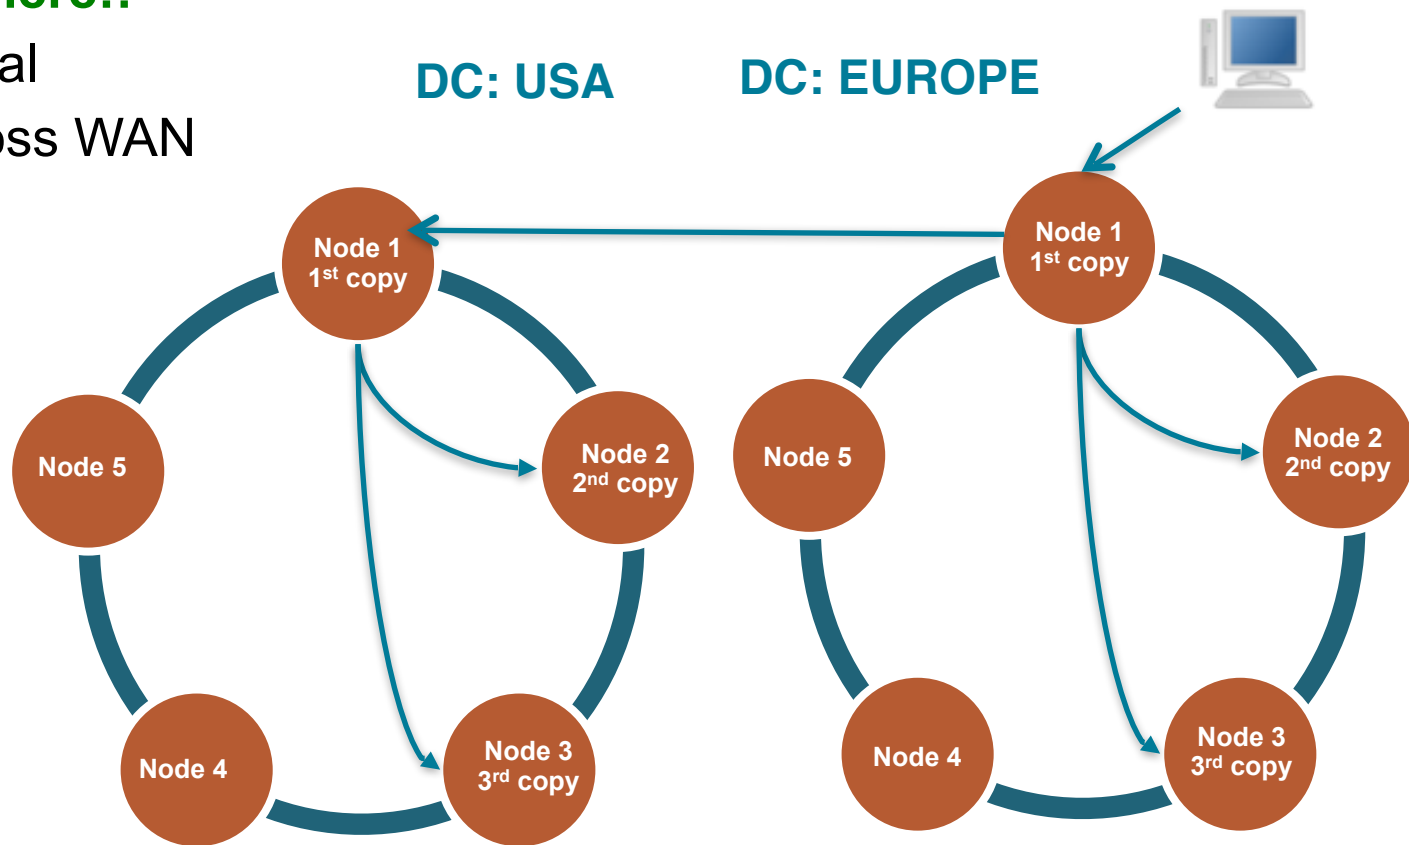
- A replication strategy determines the nodes where replicas are placed.
- All replicas are equally important; **there is no primary or master replica.**
- **Two replication strategies available:**
 - ***SimpleStrategy* (prototype only, don't use in production)**
 - ***NetworkTopologyStrategy***
- Replication strategy is set at the keyspace/schema level.



- Specified how many replicas you want in each data center (physical or virtual).
- **NetworkTopologyStrategy places replicas in the same data center by walking the ring clockwise until reaching the first node in another rack.**

Data Center Aware

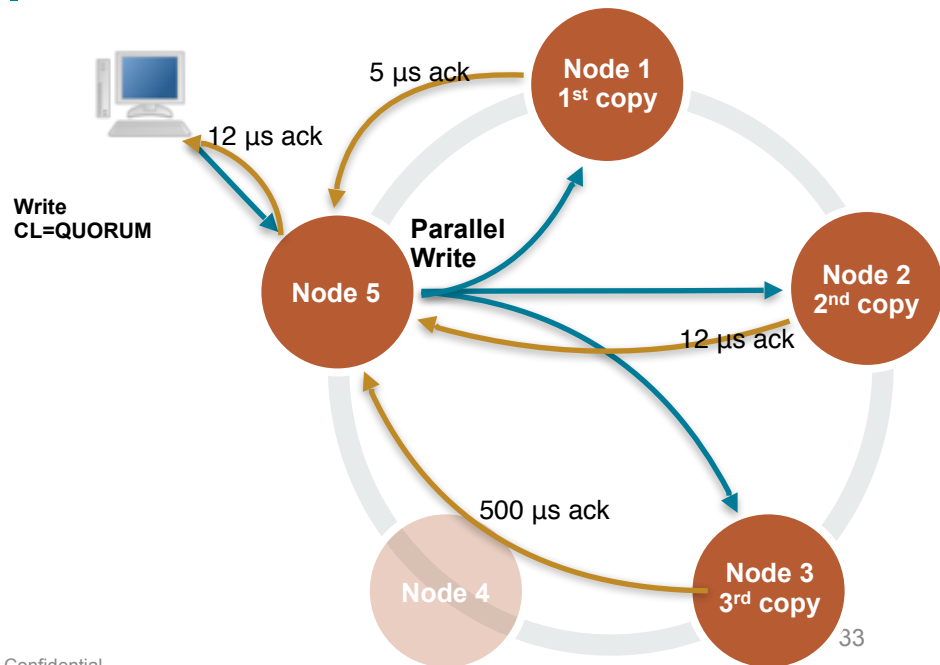
- **Active Everywhere!!**
- Client writes local
- Data syncs across WAN



Tunable Consistency

- Consistency Level (CL)
- **Client specifies per operation**
- Handles multi-data center operations

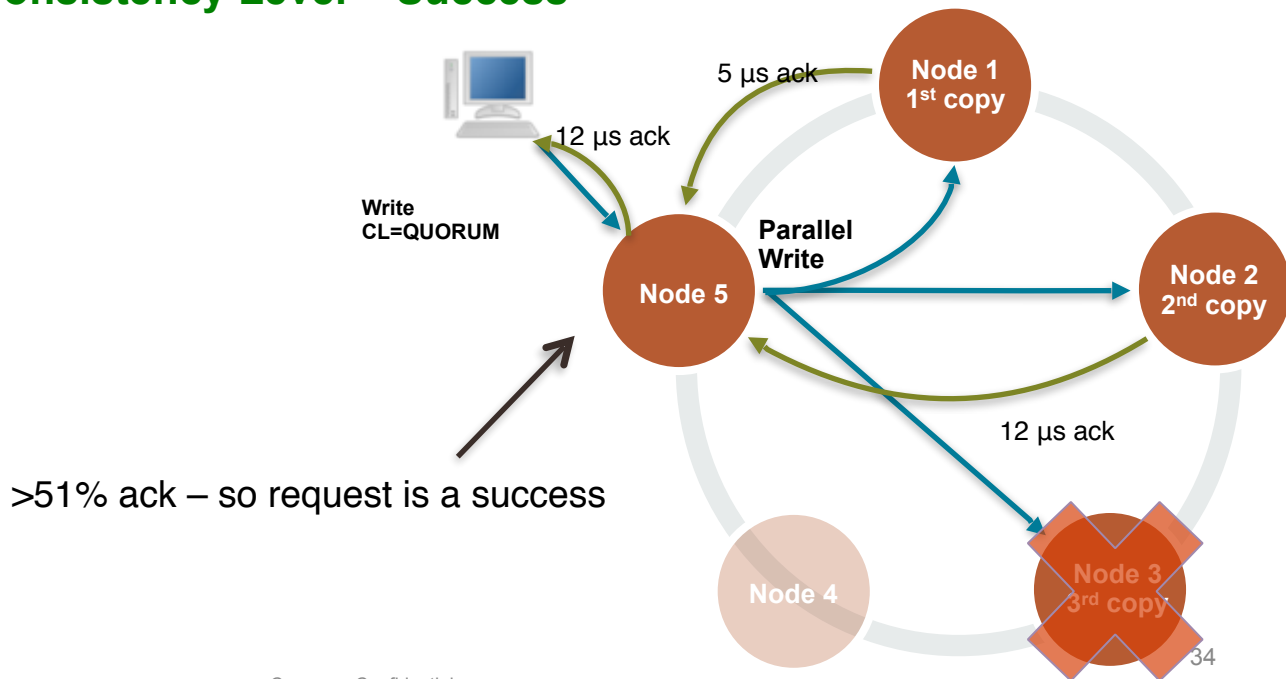
ALL = All replicas ack
 QUORUM = > 51% of replicas ack
 LOCAL_QUORUM = > 51% in local DC ack
 ONE = Only one replica acks
 Plus more.... (see docs)



Node Failure

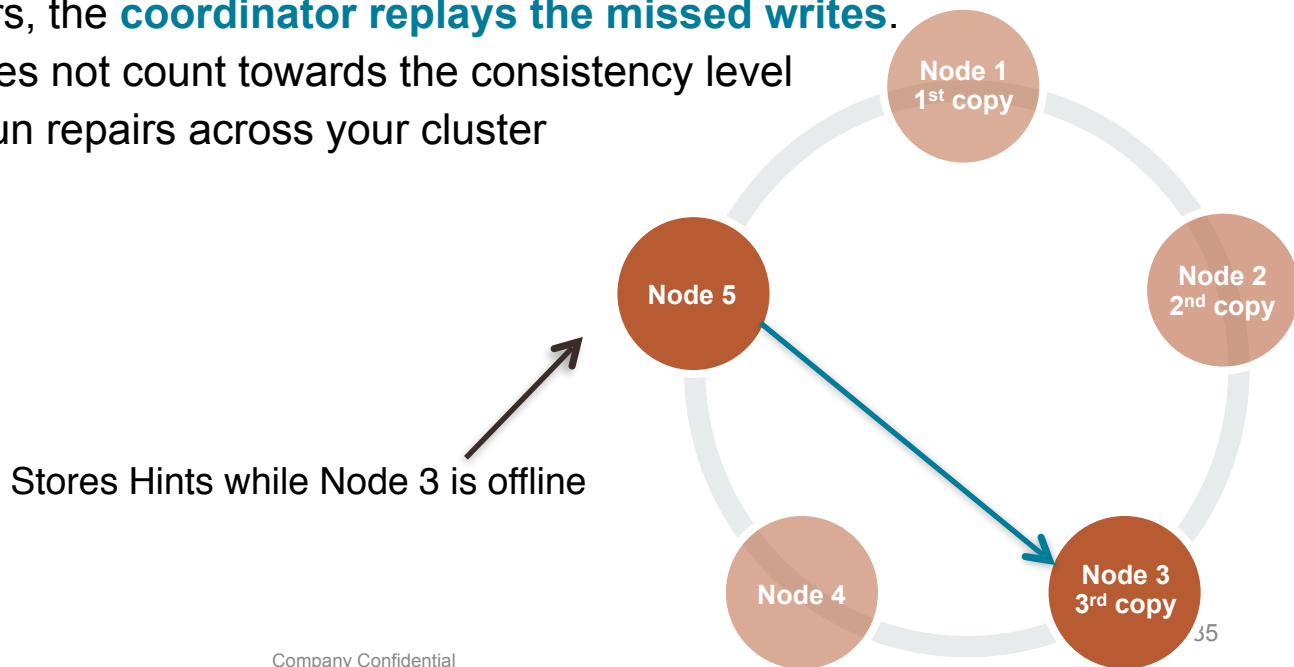
- A single node failure shouldn't bring failure.
- Replication Factor + Consistency Level = Success

- This example:
 - RF = 3
 - CL = QUORUM



Node Recovery

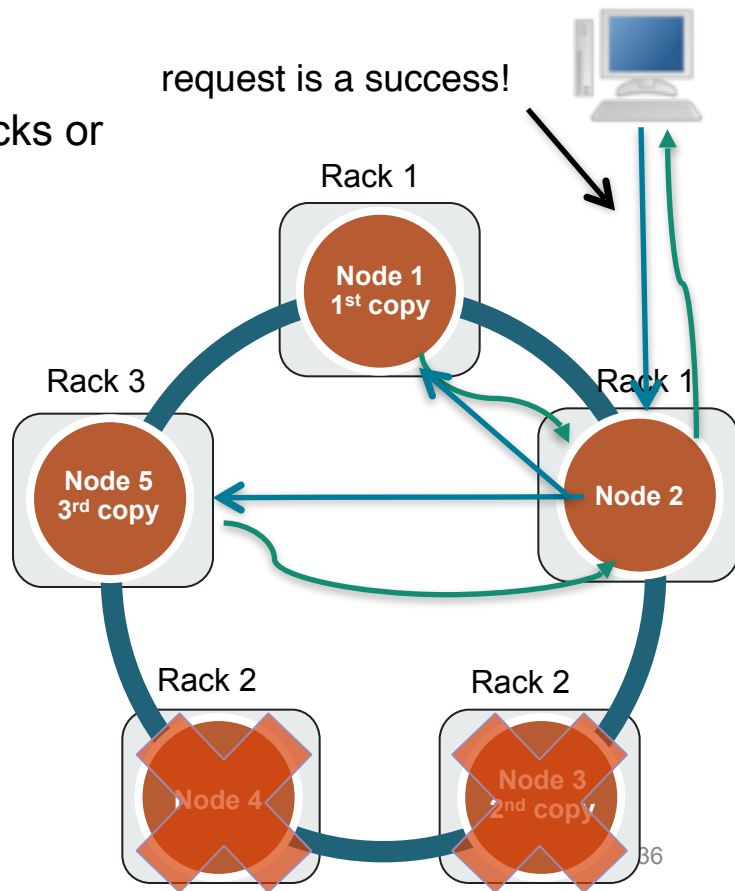
- When a write is performed and a replica node for the row is unavailable the coordinator **will store a hint locally**.
- When the node recovers, the **coordinator replays the missed writes**.
- **Note:** a hinted write does not count towards the consistency level
- **Note:** you should still run repairs across your cluster



Rack Failure

- Cassandra will place the data in as many different racks or availability zones as it can.

- This example:
 - RF = 3
 - CL = QUORUM
 - Rack 2 fails
- **Data copies still available in Node 1 and Node 5**
- **Quorum can be honored i.e. > 51% ack**



Don't be afraid of weak consistency

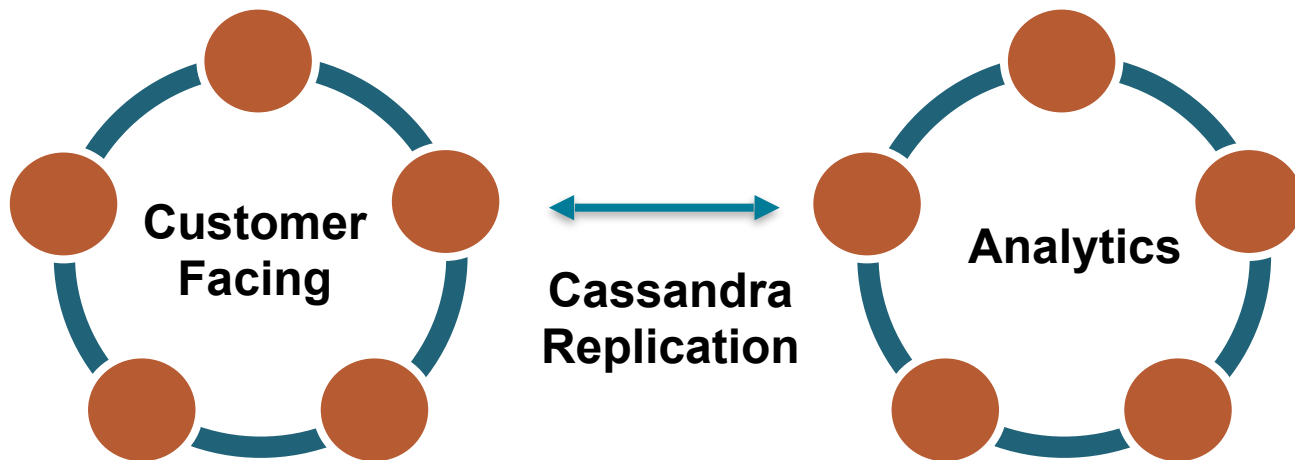


- **More tolerant to failure**
- Consistency Level of 1 is the most popular (I think)
- If you want stronger consistency go for LOCAL_QUORUM i.e. quorum is honored in the local data centre.
- If you go stronger than LOCAL_QUORUM – really understand what this means and why you are doing it.
- Remember – you can have different consistency levels for reads and writes e.g. write with CL:1, read with CL:LOCAL_QUORUM

CONSISTENCY
IS 

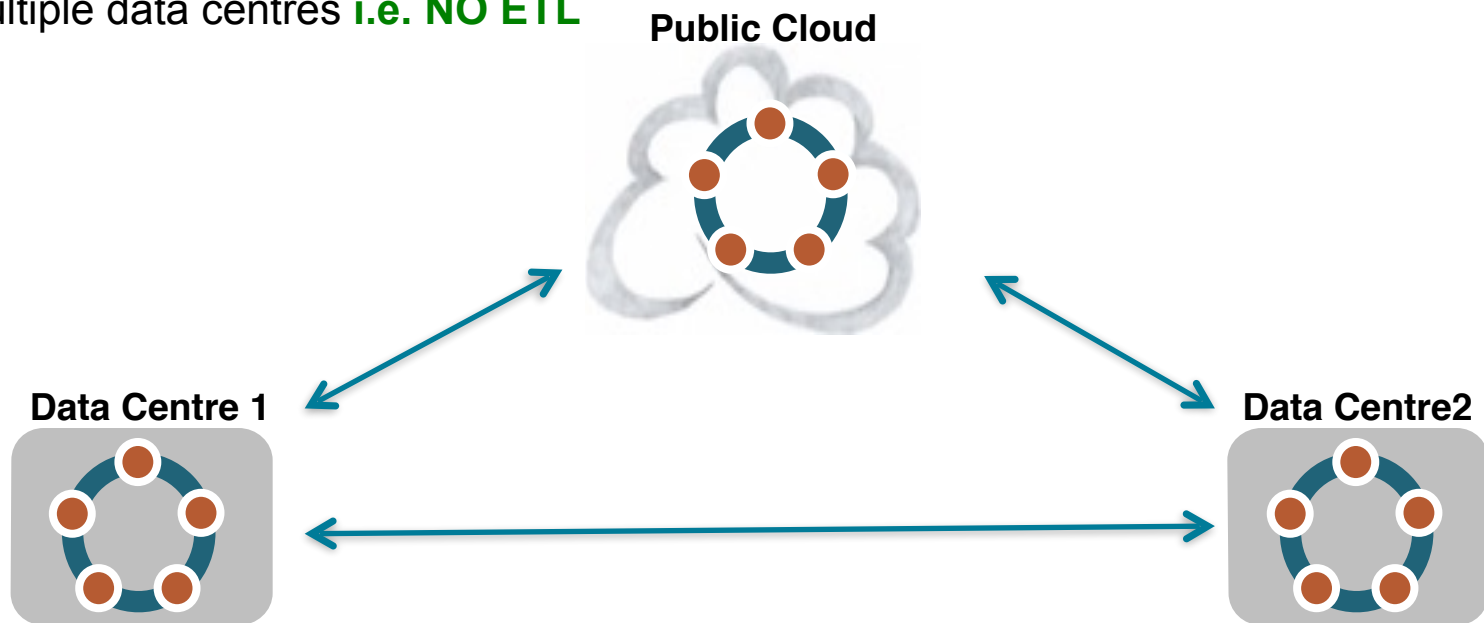
Ring fenced resources

- If you need to isolate resources for different uses, Cassandra is a great fit.
- You can create separate virtual data centres optimised as required – different workloads, hardware, availability etc..
- Cassandra will replicate the data for you automatically – **no ETL is necessary**



Hybrid Cloud

- **Cassandra is multi-data centre and cloud capable**
- Data written to any node is automatically and transparently written to all other nodes in multiple data centres **i.e. NO ETL**



Security

FEATURES



Internal Authentication
Manages login IDs and passwords inside the database



Object Permission Management
controls who has access to what and who can do what in the database



Client to Node Encryption
protects data in flight to and from a database cluster

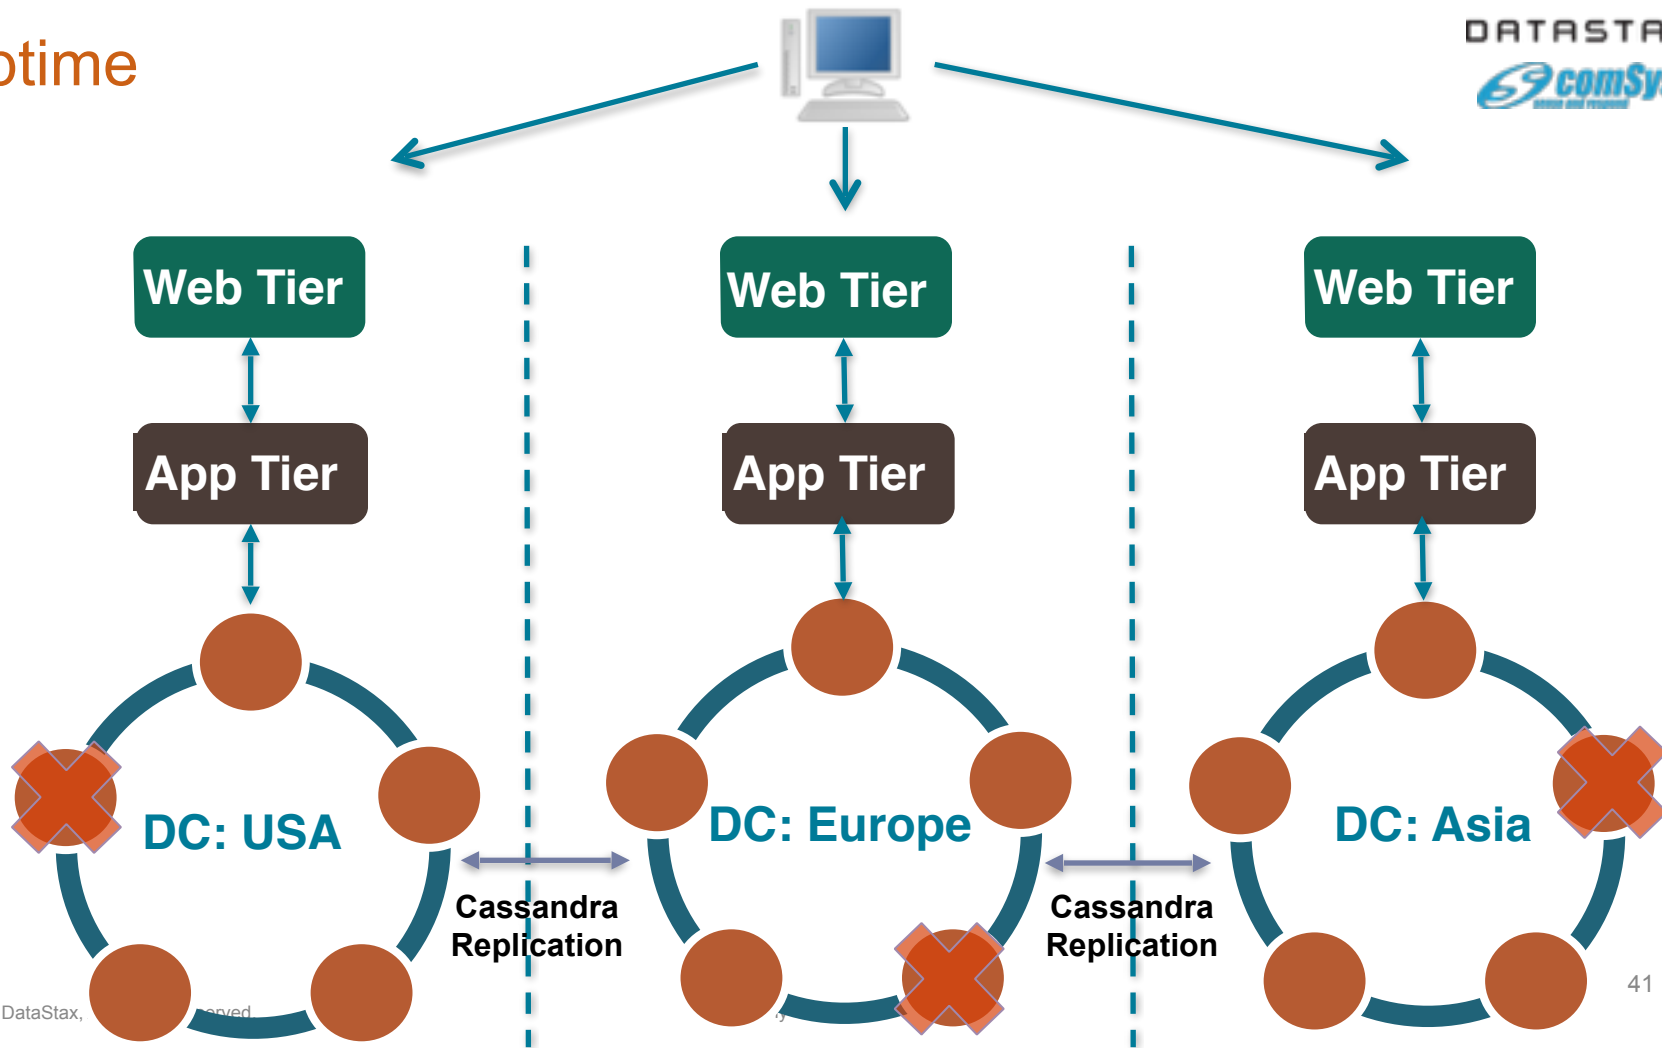
BENEFITS

- + Ensures only authorized users can access a database system using internal validation
- + Simple to implement and easy to understand
- + No learning curve from the relational world

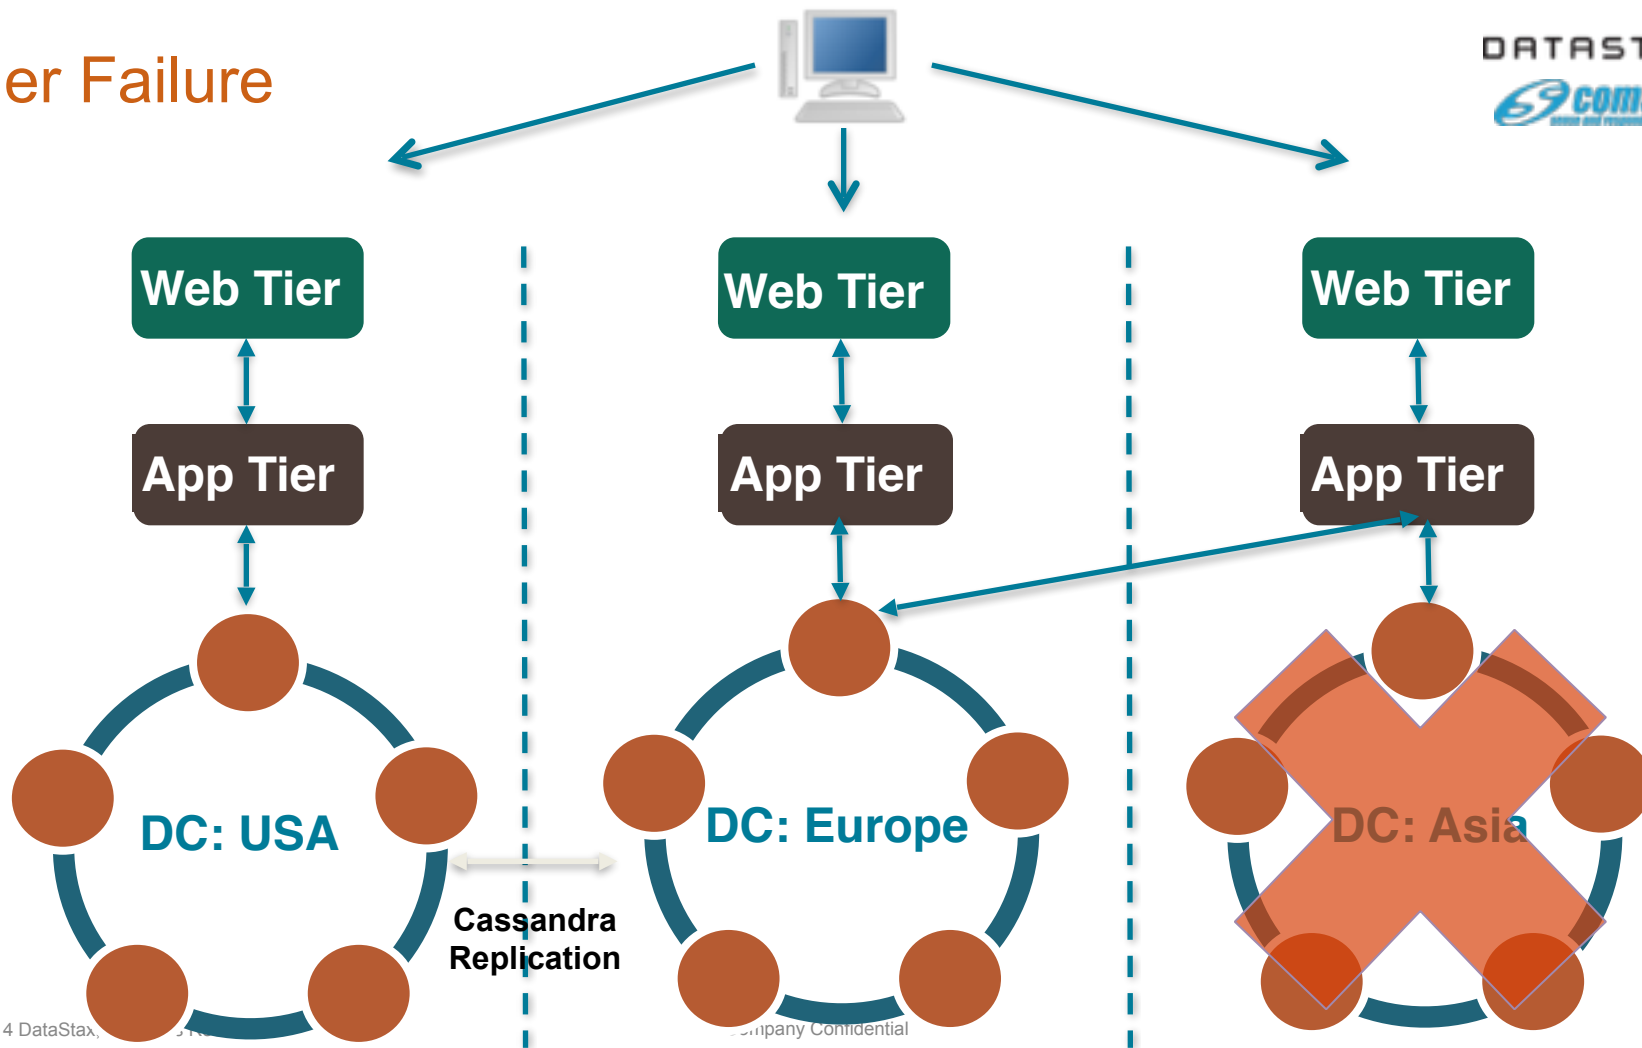
- + Provides granular based control over who can add/change/delete/read data
- + Uses familiar GRANT/REVOKE from relational systems
- + No learning curve

- + Ensures data cannot be captured/stolen in route to a server
- + Data is safe both in flight from/to a database and on the database; complete coverage is ensured

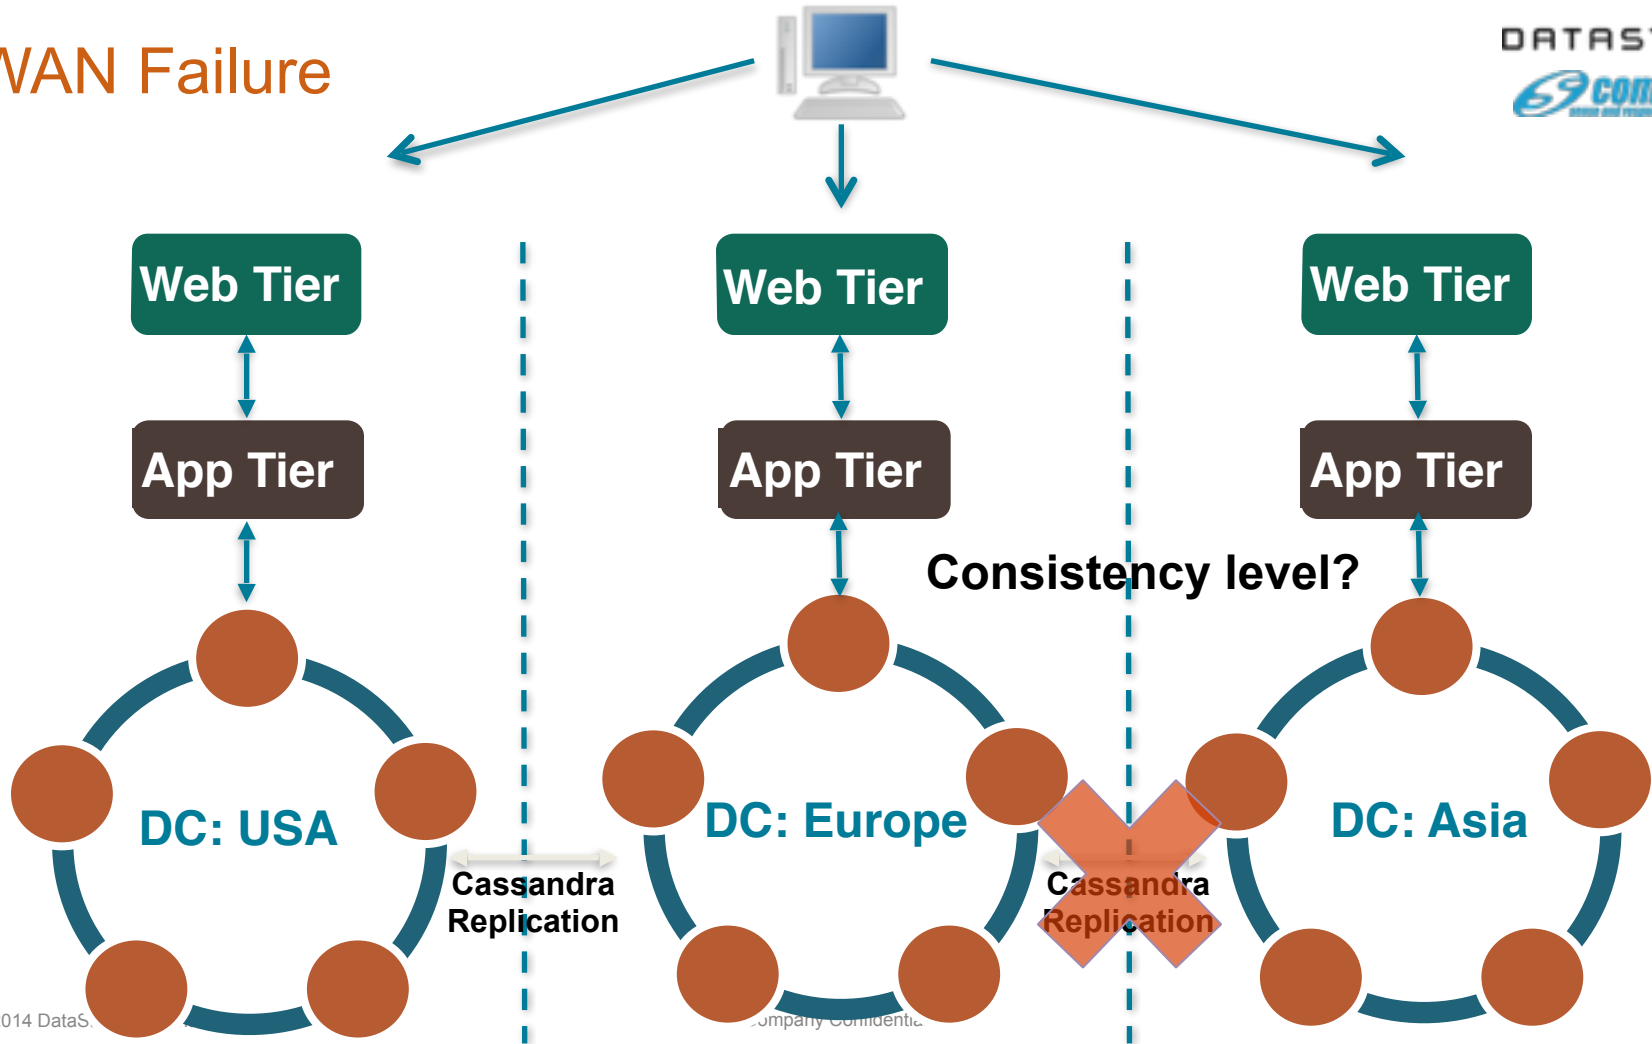
Uptime



Tier Failure



WAN Failure



How?



Data Modeling



Why Good Data Modeling is Important?

Cassandra is a highly available, highly scalable, & highly distributed database - with no single point of failure

To achieve this, **Cassandra is optimized for non-relational data models.**

- **Joins do not function well on distributed databases.**
- **Locking and transactions jam up distributed nodes**



So what do I do?

- Lots of different approaches to this...
- Basically, create multiple tables optimized for your queries.

Query Based Modeling

- The main point is:

DON'T BE AFRAID OF WRITES
DENORMALISE AND MODEL YOUR DATA AROUND YOUR
QUERIES



For more on data modeling...

- **Data modeling video series by Patrick McFadin**
- Part 1: The Data Model is Dead, Long Live the Data Model
<http://www.youtube.com/watch?v=px6U2n74q3g>
- Part 2: Become a Super Modeler
<http://www.youtube.com/watch?v=qphhxujn5Es>
- Part 3: The World's Next Top Data Model
<http://www.youtube.com/watch?v=HdJIsOZVGwM>

CQL

Cassandra Query Language

- **Cassandra Query Language**
- SQL-like language to query Cassandra
- Limited predicates. Attempts to prevent bad queries
 - but, you can still get into trouble!
- **Keyspace** – analogous to a schema.
 - Has various storage attributes.
 - The keyspace determines the RF (replication factor).
- **Table** – looks like a SQL Table.
 - A table must have a Primary Key.
 - We can fully qualify a table as <keyspace>.<table>

```

1  CREATE KEYSPACE IF NOT EXISTS cassandra;
2  USE cassandra;
3
4  CREATE TABLE IF NOT EXISTS cassandra_users (
5      name text,
6      details list<text>,
7      PRIMARY KEY (name));
8
9
10 CREATE TABLE IF NOT EXISTS cassandra_mvps (
11     userid uuid,
12     firstname varchar,
13     lastname varchar,
14     details map<text, text>,
15     PRIMARY KEY (userid));
16
17
18 //Users
19 INSERT INTO cassandra_users (name, multi_dc, details)
20 VALUES ('Netflix', true, ('http://planetcassandra.org/CompanyDetails/Netflix'));
21 INSERT INTO cassandra_users (name, details)
22 VALUES ('CEM', ('http://planetcassandra.org/blog/post/cassandra-at-cern-large-hadron-collision'));
23 INSERT INTO cassandra_users (name, details)
24 VALUES ('Metabroadcast', ('http://www.planetcassandra.org/blog/post/5-minute-c-interview--me'));
25 INSERT INTO cassandra_users (name, details)
26 VALUES ('Twitter', ('http://planetcassandra.org/CompanyDetails/Twitter'));
27
28 // Add details
29 UPDATE cassandra_users SET details = ('http://twitter.com/Netflix', 'Netflix is a high performance distributed database')
30 WHERE name = 'Netflix';
31
32
33 // MVPS
34 BEGIN BATCH
35 INSERT INTO cassandra_mvps (userid, details)
36 VALUES ('416a5ddc-00a5-40ed-adde-099da27c0c', ('twitter', 'aaronmorton'));
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999
1000
    
```

Command line interface comes with Cassandra

Launching on Linux

- `$ cqlsh [options] [host [port]]`

Launching on Windows

- `python cqlsh [options] [host [port]]`

Example

- `$ cqlsh`
- `$ cqlsh -u student -p cassandra 127.0.0.1 9160`

```

johnny@JPM-MacBook-Pro:/apps/dse/4.0/dse-4.0.2/bin$ ./cqlsh
Connected to Test Cluster at localhost:9160.
[cqlsh 4.1.1 | Cassandra 2.0.6.28 | CQL spec 3.1.1 | Thrift protocol 19.39.0]
Use HELP for help.
cqlsh> HELP

Documented shell commands:
=====
CAPTURE      COPY  DESCRIBE  EXPAND  SHOW    TRACING
CONSISTENCY  DESC  EXIT      HELP    SOURCE

CQL help topics:
=====
ALTER                CREATE_TABLE_OPTIONS  SELECT
ALTER_ADD            CREATE_TABLE_TYPES    SELECT_COLUMNFAMILY
ALTER_ALTER          CREATE_USER            SELECT_EXPR
ALTER_DROP           DELETE                 SELECT_LIMIT
ALTER_RENAME         DELETE_COLUMNS         SELECT_TABLE
ALTER_USER           DELETE_USING           SELECT_WHERE
ALTER_WITH           DELETE_WHERE           TEXT_OUTPUT
APPLY                DROP                   TIMESTAMP_INPUT
ASCII_OUTPUT         DROP_COLUMNFAMILY     TIMESTAMP_OUTPUT
BEGIN               DROP_INDEX            TRUNCATE
BLOB_INPUT          DROP_KEYSPACE         TYPES
BOOLEAN_INPUT       DROP_TABLE            UPDATE
COMPOUND_PRIMARY_KEYS DROP_USER              UPDATE_COUNTERS
CREATE              GRANT                 UPDATE_SET
CREATE_COLUMNFAMILY INSERT                UPDATE_USING
CREATE_COLUMNFAMILY_OPTIONS LIST                 UPDATE_WHERE
CREATE_COLUMNFAMILY_TYPES LIST_PERMISSIONS    USE
CREATE_INDEX        LIST_USERS            UUID_INPUT
CREATE_KEYSPACE     PERMISSIONS
CREATE_TABLE        REVOKE

cqlsh>

```

- Usual statements
- CREATE / DROP / ALTER TABLE
- SELECT / INSERT / UPDATE
- **Plus many, many more!**

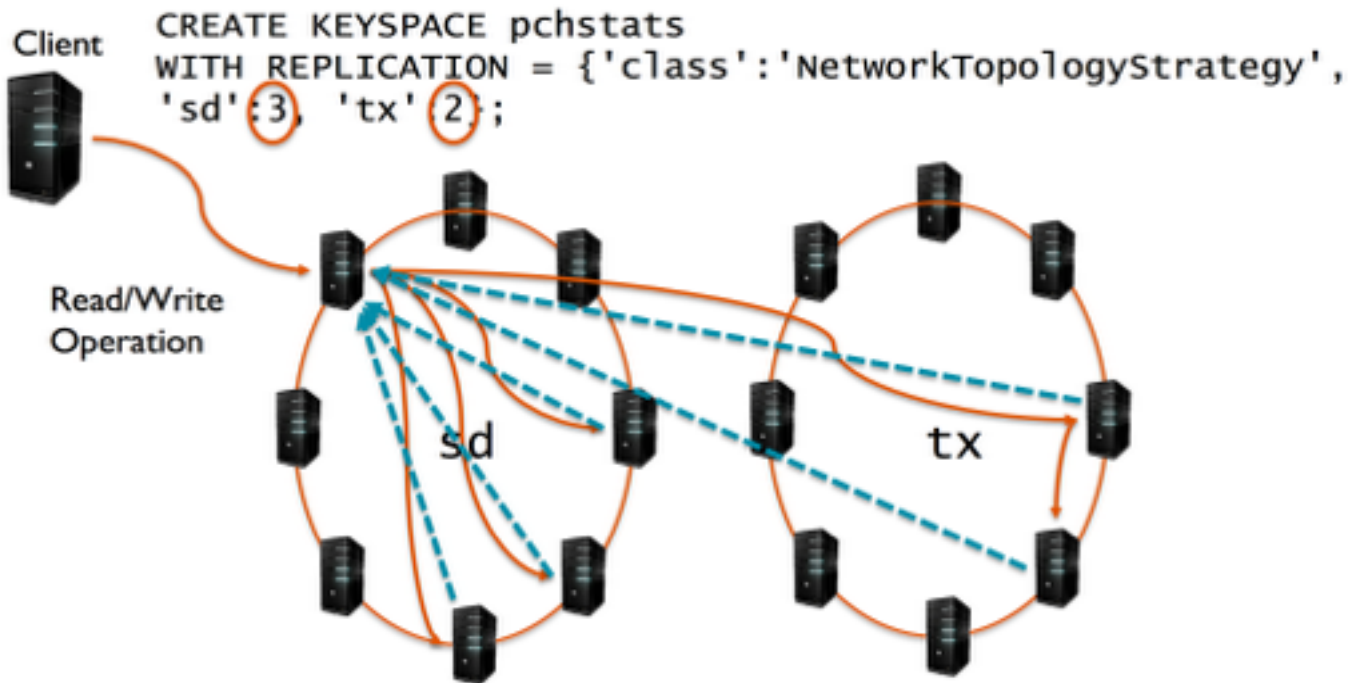
www.datastax.com/docs

What is a keyspace?

- **Keyspace or schema is a top-level namespace**
 - All data objects (e.g., tables) must belong to some keyspace
 - Defines how data is replicated on nodes
 - Keyspace per application is a good idea
- **Replica placement strategy**
 - **SimpleStrategy (prototyping)**
 - **NetworkTopologyStrategy (production)**



Creating a keyspace - Multiple Data Centre



Use and Drop a keyspace

- To work with data objects (e.g., tables) in a keyspace:
- `USE pchstats;`
- To delete a keyspace and all internal data objects
- `DROP KEYSPACE pchstats;`

If you drop a keyspace or table, Cassandra will back it up!
You can disable this via `auto_snapshot` in the `cassandra.yaml`



Creating a table

```
CREATE TABLE cities (  
    city_name    varchar,  
    elevation    int,  
    population   int,  
    latitude     float,  
    longitude    float,  
    PRIMARY KEY (city_name)  
);
```

- We can visualize it this way:

city_name	elevation	population	latitude	longitude
Springfield	978	60608	39°55'37"	83°48'15" W

- Here, city_name is the partition key
- In this example, the partition key = primary key

The Primary Key

- The key uniquely identifies a row.
- A composite primary key consists of:
 - A **partition key**
 - One or more **clustering columns**
- e.g. `PRIMARY KEY (partition key, cluster columns, ...)`
- The **partition key** determines on which node the partition resides
- Data is **ordered** in **cluster column** order within the partition

Compound Primary Key

```
CREATE TABLE sporty_league (  
    team_name    varchar,  
    player_name  varchar,  
    jersey       int,  
    PRIMARY KEY (team_name, player_name)  
);
```

Clustering Column

Partition Key

team_name Springers	Adler 86	Bélanger 13	Foote 99
team_name Mighty Mutts	Buddy 32	Lucky 7	
team_name Peppers	Aaron 17	Baker 62	Cabrera 25

Not ordered

Ordered by player_name

Simple Select

This is a bad query



```
SELECT * FROM sporty_league;
```

Partition Key

team_name	player_name	jersey
Peppers	Aaron	17
Peppers	Baker	62
Peppers	Cabrera	25
Springers	Adler	86
Springers	Bélanger	13
Springers	Footé	99
Mighty Mutts	Buddy	32
Mighty Mutts	Lucky	7

Clustering Column
ordered

- More than a few rows can be slow. (Limited to 10,000 rows by default)
- Use **LIMIT** keyword to choose fewer or more rows

Select on Partition Key and Cluster Columns

```
SELECT * FROM sporty_league WHERE team_name = 'Mighty Mutts';
```

team_name	player_name	jersey
Mighty Mutts	Buddy	32
Mighty Mutts	Lucky	7

```
SELECT * FROM sporty_league WHERE team_name = 'Mighty Mutts'  
and player_name = 'Lucky';
```

team_name	player_name	jersey
Mighty Mutts	Lucky	7

Composite Partition Key

Partition Key

```
CREATE TABLE cities (  
    city_name    varchar,  
    state        varchar  
    PRIMARY KEY ((city_name, state))  
);
```

(city_name, state)	elevation	latitude	longitude	population
Springfield, OH	978	39.926	-83.804	60608
(city_name, state)	elevation	latitude	longitude	population
Springfield, IL	39.783	39.783	-89.650	116250

CQL Predicates

- On the **partition key**: = and IN
- On the **cluster columns**: <, <=, =, >=, >, IN

Performance Considerations

- The best queries are in a single partition.
i.e. **WHERE partition key = <something>**
- Each new partition requires a new disk seek.
- Queries that span multiple partitions are **s-l-o-w**
- Queries that span multiple cluster columns are **fast**



Ordering

- Partition keys are **not** ordered, but the cluster columns are.
- However, **you can only order by a column if it's a cluster column.**
- Data will returned by default in the order of the clustering column.
- You can also use the ORDER BY keyword – **but only on the clustering column!**

```
SELECT * FROM sporty_league  
WHERE team_name = 'Mighty Mutts'  
ORDER BY player_name DESC;
```



Secondary Indexes

If we want to do a query on a column that is not part of your PK, you can create an index:

```
CREATE INDEX ON <table>(<column>);
```

Then you can do a select:

```
SELECT * FROM product WHERE type= 'PC';
```

Suitable for low cardinality data

Much more efficient to model your data around the query



Watch out for global indexes in Cassandra 3.0 (CASSANDRA-6477)

```
INSERT INTO sporty_league (team_name, player_name, jersey)
VALUES ('Mighty Mutts', 'Felix', 90);
```

```
UPDATE sporty_league SET jersey = 77
WHERE team_name = 'Mighty Mutts' AND player_name = 'Felix';
```

Inserts a row into a table

- Must specify columns to insert values into
- **Primary key columns are mandatory (identify the row)**
- Other columns do not have to have values and **Non-existent 'values' do not take up space**

Inserts are atomic

- All values of a row are inserted or none

Inserts are isolated

- Two inserts with the same values in primary key columns will not interfere – executed one after another

What is an upsert?

- **UPDATE + inSERT**
 - Both UPDATE and INSERT are write operations
 - No reading before writing
- **Term “upsert” denotes the following behavior**
 - INSERT updates or overwrites an existing row
 - When inserting a row in a table that already has another row with the same values in primary key columns
 - UPDATE inserts a new row
 - When a to-be-updated row, identified by values in primary key columns, does not exist
- **Upserts are legal and do not result in error or warning messages**

How to avoid UPSERTS

Guarantee that your primary keys are unique from one another

- Use an appropriate natural key based on your data

- Use a surrogate key for partition key

Risks with natural keys

- Depending on the type of natural key that is used, there may still be an increased risk of UPSERTs

- Changing the datum used for a Natural Key requires a lot of overhead.

So why not use a sequence to generate a surrogate key?

- You cant – **Cassandra doesn't provide sequences!**

What? No sequences!

- **Sequences are a handy feature in RDMBS for auto-creation of IDs for you data.**
 - Guaranteed unique
 - E.g. `INSERT INTO user (id, firstName, LastName) VALUES (seq.nextVal(), 'Ted', 'Codd')`
- **Cassandra has no sequences!**
 - Extremely difficult in a masterless distributed system
 - **Requires a lock (perf killer)**
- What to do?
 - **Use part of the data to create a unique key**
 - **Use a UUID or TIMEUUID**



- Universal Unique ID
- 128 bit number represented in character form e.g.
99051fe9-6a9c-46c2-b949-38ef78858dd0
- Easily generated on the client
- Version 1 has a timestamp component (TIMEUUID)
- Version 4 has no timestamp component
 - Faster to generate

TIMEUUID data type supports Version 1 UUIDs

Generated using time (60 bits), a clock sequence number (14 bits), and MAC address (48 bits)

- **CQL function 'now()' generates a new TIMEUUID**

Time can be extracted from TIMEUUID

- **CQL function dateOf() extracts the timestamp as a date**

TIMEUUID values in clustering columns or in column names are ordered based on time

- **DESC order on TIMEUUID lists most recent data first**

Time-to-Live (TTL)

TTL a row:

```
INSERT INTO users (id, first, last) VALUES ('abc123', 'catherine', 'cachart')  
USING TTL 3600; // Expires data in one hour
```

TTL a column:

```
UPDATE users USING TTL 30 SET last = 'miller' WHERE id = 'abc123'
```

- TTL in seconds
- **Can also set default TTL at a table level**
- Expired columns/values automatically deleted
- With no TTL specified, columns/values never expire
- TTL is useful for automatic deletion
- Re-inserting the same row before it expires will overwrite TTL

Collection Data Type

- CQL supports having columns that contain collections of data.
- The collection types include:

- **Set, List and Map.**

```
CREATE TABLE collections_example (
    id int PRIMARY KEY,
    set_example set<text>,
    list_example list<text>,
    map_example map<int, text>
);
```

- These data types are intended to support the type of one-to-many relationships that can be modeled in a relational DB e.g. a user has many email addresses.
- **Some performance considerations around collections.**
 - Often more efficient to denormalise further rather than use collections if intending to store lots of data.
 - **Favor sets over list**

You now have collection indexing in Cassandra 2.1



Lightweight Transactions (LWT)

Why?

- Solve a class of race conditions in Cassandra that you would otherwise need to install an external locking manager to solve.

Syntax:

```
INSERT INTO customer_account (customerID, customer_email)
VALUES ('Johnny', 'jmilller@datastax.com')
IF NOT EXISTS;
```

```
UPDATE customer_account
SET customer_email='jmilller@datastax.com'
IF customer_email='jmilller@datastax.com';
```

Example Use Case:

- Registering a user



Not Will Ferrell @itsWillyFerrell · Apr 5

In about 20 years, the hardest thing our kids will have to do is find a username that isn't taken.

Race Condition

```
SELECT name
FROM users
WHERE username = 'johnny';

(0 rows)
```

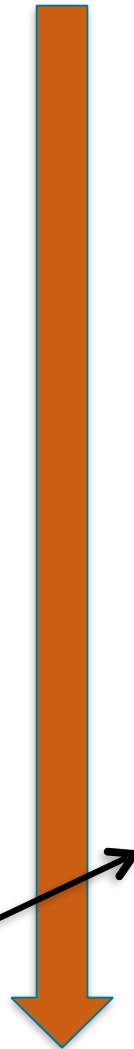
```
INSERT INTO users
(username, name, email,
password, created_date)
VALUES ('johnny',
'Johnny Miller',
['jmiller@datastax.com'],
'ba27e03fd9...',
'2011-06-20 13:50:00');
```

```
SELECT name
FROM users
WHERE username = 'johnny';

(0 rows)
```

```
INSERT INTO users
(username, name, email,
password, created_date)
VALUES ('johnny',
'Johnny Miller',
['jmiller@datastax.com'],
'ea24e13ad9...',
'2011-06-20 13:50:01');
```

This one wins!



Lightweight Transactions (LWT)

```
INSERT INTO users
(username, name, email,
password, created_date)
VALUES ('johnny',
'Johnny Miller',
['jmilller@datastax.com'],
'ba27e03fd9...',
'2011-06-20 13:50:00')
IF NOT EXISTS;
```

[applied]

True

```
INSERT INTO users
(username, name, email,
password, created_date)
VALUES ('johnny',
'Johnny Miller',
['jmilller@datastax.com'],
'ea24e13ad9...',
'2011-06-20 13:50:01')
IF NOT EXISTS;
```

[applied]	username	created_date	name
False	johnny	2011-06-20 ...	Johnny Miller

Lightweight Transactions (LWT)

Consequences of Lightweight Transactions

4 round trips vs. 1 for normal updates (uses Paxos algorithm)

Operations are done on a per-partition basis

Will be going across data centres to obtain consensus (unless you use LOCAL_SERIAL consistency)

Cassandra user will need read and write access i.e. you get back the row!



Great for 1% your app, but eventual consistency is still your friend!

Batch Statements

```
BEGIN BATCH
```

```
  INSERT INTO users (userID, password, name) VALUES ('user2', 'ch@ngem3b', 'second user')
```

```
  UPDATE users SET password = 'ps22dhds' WHERE userID = 'user2'
```

```
  INSERT INTO users (userID, password) VALUES ('user3', 'ch@ngem3c')
```

```
  DELETE name FROM users WHERE userID = 'user2'
```

```
APPLY BATCH;
```

BATCH statement combines multiple INSERT, UPDATE, and DELETE statements into a single logical operation

Atomic operation

If any statement in the batch succeeds, all will

No batch isolation

Other “transactions” can read and write data being affected by a partially executed batch

User Defined Types (UDT's)

```
CREATE TYPE address (  
    street text,  
    city text,  
    zip_code int,  
    phones set<text>  
)
```

```
CREATE TABLE users (  
    id uuid PRIMARY KEY,  
    name text,  
    addresses map<text, frozen <address>>  
)
```

```
SELECT id, name, addresses.city, addresses.phones FROM users;
```

id	name	addresses.city	addresses.phones
63bf691f	johnny	London	{ '0201234567', '0796622222' }

Lets store some JSON

```
{
  "productId": 2,
  "name": "Kitchen Table",
  "price": 249.99,
  "description": "Rectangular table with oak finish",
  "dimensions": {
    "units": "inches",
    "length": 50.0,
    "width": 66.0,
    "height": 32
  },
  "categories": {
    {
      "category": "Home Furnishings" {
        "catalogPage": 45,
        "url": "/home/furnishings"
      },
      {
        "category": "Kitchen Furnishings" {
          "catalogPage": 108,
          "url": "/kitchen/furnishings"
        }
      }
    }
  }
}
```

Lets store some JSON

```
{
  "productId": 2,
  "name": "Kitchen Table",
  "price": 249.99,
  "description": "Rectangular table with oak finish",
  "dimensions": {
    "units": "inches",
    "length": 50.0,
    "width": 66.0,
    "height": 32
  },
  "categories": {
    {
      "category": "Home Furnishings" {
        "catalogPage": 45,
        "url": "/home/furnishings"
      },
    {
      "category": "Kitchen Furnishings" {
        "catalogPage": 108,
        "url": "/kitchen/furnishings"
      }
    }
  }
}
```

```
CREATE TYPE dimensions (
  units text,
  length float,
  width float,
  height float
)
```

Lets store some JSON

```
{
  "productId": 2,
  "name": "Kitchen Table",
  "price": 249.99,
  "description": "Rectangular table with oak finish",
  "dimensions": {
    "units": "inches",
    "length": 50.0,
    "width": 66.0,
    "height": 32
  },
  "categories": {
    {
      "category": "Home Furnishings" {
        "catalogPage": 45,
        "url": "/home/furnishings"
      },
      {
        "category": "Kitchen Furnishings" {
          "catalogPage": 108,
          "url": "/kitchen/furnishings"
        }
      }
    }
  }
}
```

```
CREATE TYPE dimensions (
  units text,
  length float,
  width float,
  height float
)
```

```
CREATE TYPE category (
  catalogPage int,
  url text
);
```

Lets store some JSON

```
{
  "productId": 2,
  "name": "Kitchen Table",
  "price": 249.99,
  "description": "Rectangular table with oak finish",
  "dimensions": {
    "units": "inches",
    "length": 50.0,
    "width": 66.0,
    "height": 32
  },
  "categories": {
    {
      "category": "Home Furnishings" {
        "catalogPage": 45,
        "url": "/home/furnishings"
      },
      {
        "category": "Kitchen Furnishings" {
          "catalogPage": 108,
          "url": "/kitchen/furnishings"
        }
      }
    }
  }
}
```

```
CREATE TYPE dimensions (
  units text,
  length float,
  width float,
  height float
)
```

```
CREATE TYPE category (
  catalogPage int,
  url text
);
```

```
CREATE TABLE product (
  productId int,
  name text,
  price float,
  description text,
  dimensions frozen <dimensions>,
  categories map <text,frozen <category>>,
  PRIMARY KEY (productId)
);
```

What's frozen mean?



- This applies to User Defined Types, Tuples and nested collections.
- Frozen types **are serialized as a single value in Cassandra's storage engine**, whereas **non-frozen types are stored in a form that allows updates to individual subfields**.
- When using the frozen keyword, you cannot update parts of a user-defined type value. **The entire value must be overwritten.**
- Cassandra treats the value of a frozen type like a blob.
- **Cassandra 3.0 will support non-frozen**. As of 2.1, we only support frozen UDTs.
- Helps us avoid tech debt!



Lets store some JSON

```
INSERT INTO product (productId, name, price,
description, dimensions, categories)
VALUES (2, 'Kitchen Table', 249.99, 'Rectangular
table with oak finish',
```

```
{
  units: 'inches',
  length: 50.0,
  width: 66.0,
  height: 32
```



dimensions frozen <dimensions>

```
},
{
  'Home Furnishings': {
    catalogPage: 45,
    url: '/home/furnishings'
  },
  'Kitchen Furnishings': {
    catalogPage: 108,
    url: '/kitchen/furnishings'
  }
}
```



categories map <text, frozen <category>>

```
}
);
```

Query Tracing

- You can turn on tracing on or off for queries with **the TRACING ON | OFF** command.
- Helps you understand what Cassandra is doing and identify any performance problems.



```
cqlsh:ecow> SELECT vendor, order_id, user_id, quantity, total_cost, product_id, product_name, order_timestamp FROM order_by_vendor WHERE vendor='Yoncos BBQ & Grill Franchising' AND bucket = 1;
```

vendor	order_id	user_id	quantity	total_cost	product_id	product_name	order_timestamp
Yoncos BBQ & Grill Franchising	0231	01949	0	110.68	P1632	Salsa - Cranberry Grapefruit	2013-08-08 23:02:10+0000

Tracing session: 9d516d98-d743-11e3-b6a1-116640baf3d0

activity	timestamp	source	source_elapsed
		execute_cql3_query	0
Parsing SELECT vendor, order_id, user_id, quantity, total_cost, product_id, product_name, order_timestamp FROM order_by_vendor WHERE vendor='Yoncos BBQ & Grill Franchising' AND bucket = 1 LIMIT 10000;	08:29:07,031	192.168.104.176	140
	08:29:07,031	192.168.104.176	350
	08:29:07,032	192.168.104.176	1005
	08:29:07,032	192.168.104.176	1120
	08:29:07,032	192.168.104.176	1155
	08:29:07,032	192.168.104.176	1271
	08:29:07,032	192.168.104.176	1354
	08:29:07,032	192.168.104.176	1600

CQL supports creating users and granting them access to tables etc..

You need to enable authentication in the cassandra.yaml config file.

You can create, alter, drop and list users

You can then GRANT permissions to users accordingly – ALTER, AUTHORIZE, DROP, MODIFY, SELECT.



Plus lots, lots more you can do with CQL

- Have a look at our documentation:
<http://www.datastax.com/documentation>
- Technical Blog
<http://www.datastax.com/dev/blog>

DataStax Java Driver



How to get it?

- <https://github.com/datastax/java-driver>

```
<dependency>  
  <groupId>com.datastax.cassandra</groupId>  
  <artifactId>cassandra-driver-core</artifactId>  
  <version>2.1.0</version>  
</dependency>
```

- The Java client driver 2.1 (branch 2.1) is **compatible** with Apache Cassandra 1.2, 2.0 and 2.1.
- If you try to use a feature that's not in the version of Cassandra you are connecting to, you will get an **UnsupportedFeatureException**

- Traditionally, Cassandra clients (Hector, Astynax¹ etc..) were developed using Thrift
- With Cassandra 1.2 (Jan 2013) and the introduction of CQL3 and **the CQL native protocol** a new easier way of using Cassandra was introduced.
- Why?
 - Easier to develop and model
 - Best practices for building modern distributed applications
 - Integrated tools and experience
 - Enable Cassandra to evolve easier and support new features

¹Astynax is being updated to include the native driver: <https://github.com/Netflix/astyanax/wiki/Astyanax-over-Java-Driver>

Thrift post-Cassandra 2.1



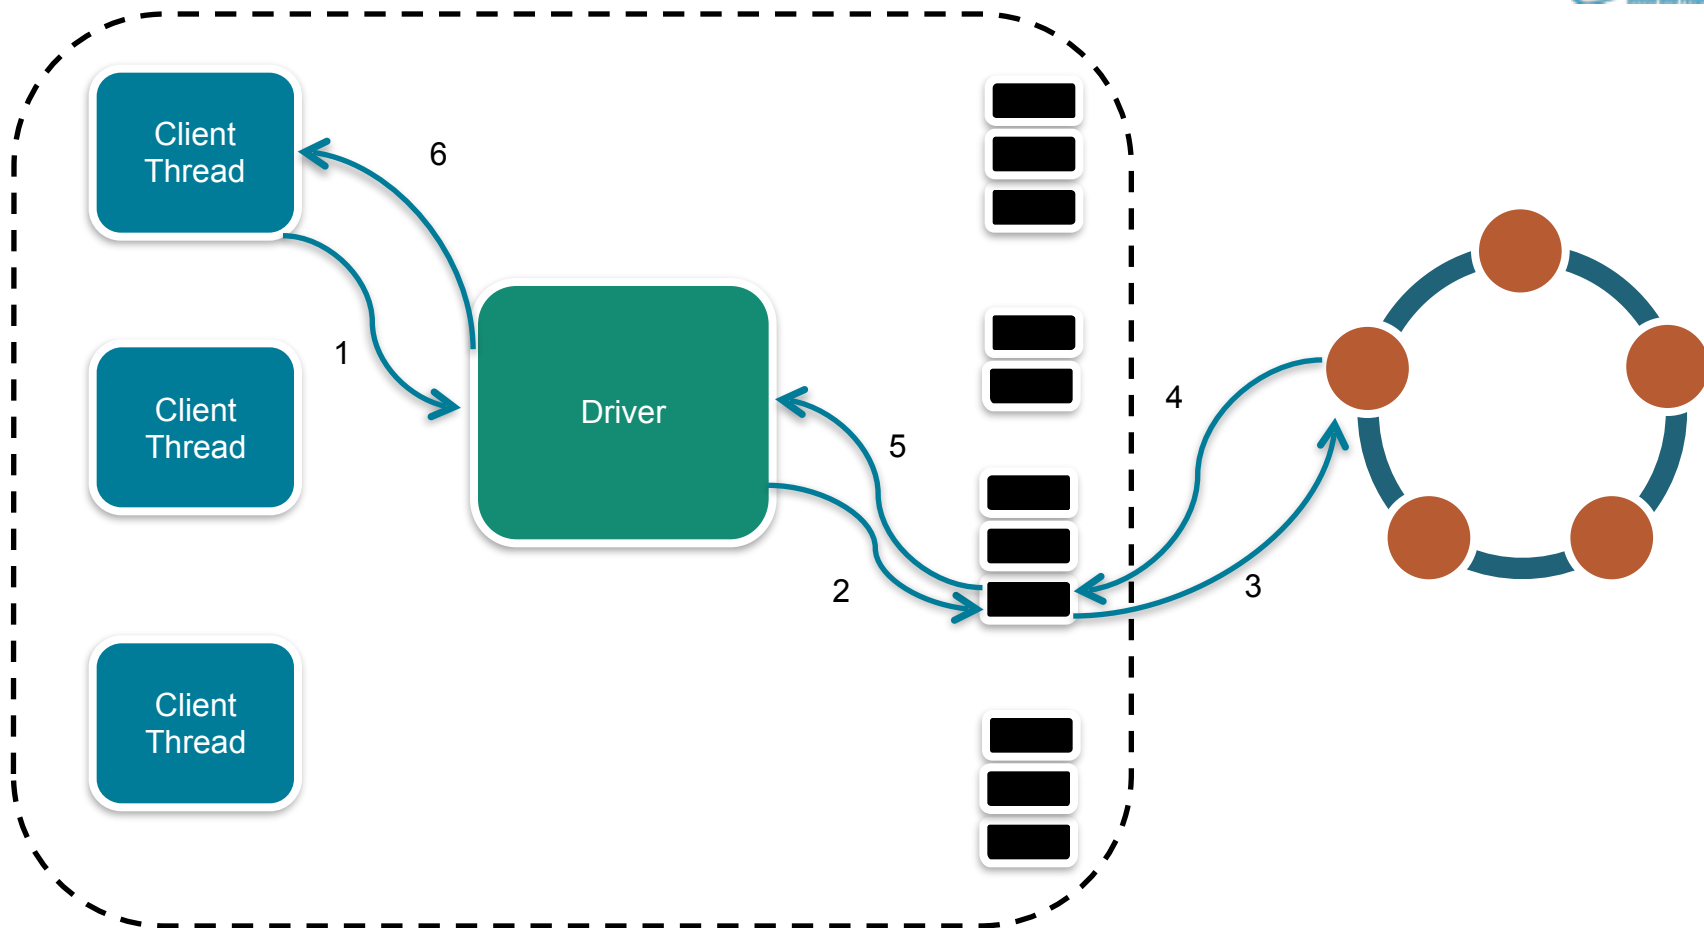
There is no more work on Thrift post 2.1.0

- <http://bit.ly/freeze thrift>

Will retain it for backwards compatibility, but no new features or changes to the Thrift API after 2.1.0

- *“**CQL3 is almost two years old** now and has proved to be the **better API** that Cassandra needed. CQL drivers have caught up with and passed the Thrift ones in terms of **features, performance, and usability**. CQL is **easier to learn and more productive** than Thrift.”*
- - Jonathan Ellis, Apache Chair, Cassandra

Asynchronous Architecture



Connect and Write

```
Cluster cluster = Cluster.builder()
    .addContactPoints("10.158.02.40", "10.158.02.44")
    .build();
```

```
Session session = cluster.connect("akeyspace");
```

```
session.execute(
    "INSERT INTO user (username, password) "
    + "VALUES('johnny', 'password1234') "
);
```

Note: Clusters and Sessions should be long-lived and re-used.

Read from a table

```
ResultSet rs = session.execute("SELECT * FROM user");
```

```
List<Row> rows = rs.all();
```

```
for (Row row : rows) {  
    String userName = row.getString("username");  
    String password = row.getString("password");  
}
```

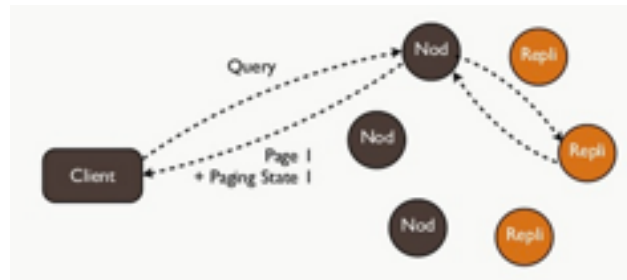
Paging

This is great!

Historically difficult to get huge result sets out of Cassandra. It has generally been necessary to explicitly enumerate your row keys in reasonably small batches (1000 rows or so per batch would be common).

This feature now allows you to get huge result sets (including “select * from table), and have the server automatically page the results, while the client is just able to trivially iterate over the entire result set.

Makes data exploration much easier.



Asynchronous Requests

```
ResultSetFuture future = session.executeAsync("SELECT * FROM user");

for (Row row : future.get()) {
    String userName = row.getString("username");
    String password = row.getString("password");
}
```

Note: The future returned implements Guava's `ListenableFuture` interface. This means you can use all Guava's Futures¹ methods!

¹<http://docs.guava-libraries.googlecode.com/git/javadoc/com/google/common/util/concurrent/Futures.html>

Read with Callbacks

```
final ResultSetFuture future = session.executeAsync("SELECT * FROM user");
Futures.addCallback(future, new FutureCallback<ResultSet>() {
    @Override
    public void onFailure(Throwable t) {
        log.error("Problem encountered", t);
        //Do something...
    }

    @Override
    public void onSuccess(ResultSet r) {
        log.debug("Success!!");
        //Do something...
    }
})
```

Parallelize Calls

```
int queryCount = 99;

List<ResultSetFuture> futures = new ArrayList<ResultSetFuture>();

for (int i=0; i<queryCount; i++) {
    futures.add(
        session.executeAsync("SELECT * FROM user "
                               + "WHERE username = '"+i+"'"));
}

for(ResultSetFuture future : futures) {
    for (Row row : future.getUninterruptibly()) {
        //do something
    }
}
```

- If you need to do a lot of work, **it's often better to make many small queries concurrently than to make one big query.**
 - **executeAsync and Futures – makes this really easy!**
 - Big queries can put a high load on one coordinator
 - Big queries can skew your 99th percentile latencies for other queries
 - If one small query fails you can easily retry, if a big query than you have to retry the whole thing



```
Query query = QueryBuilder
    .select()
    .all()
    .from("akeyspace", "user")
    .where(eq("username", "johnny"));

query.setConsistencyLevel(ConsistencyLevel.ONE);

ResultSet rs = session.execute(query);
```

Multi Data Center Load Balancing

Local nodes are queried first, if none are available the request will be sent to a remote data center

```
Cluster cluster = Cluster.builder()  
    .addContactPoints("10.158.02.40", "10.158.02.44")  
    .withLoadBalancingPolicy(  
        new DCAwareRoundRobinPolicy("DC1"))  
    .build();
```

Name of the local DC

Nodes that own a replica of the data being read or written by the query will be contacted first

```
Cluster cluster = Cluster.builder()  
    .addContactPoints("10.158.02.40", "10.158.02.44")  
    .withLoadBalancingPolicy(  
        new TokenAwarePolicy(  
            new DCAwareRoundRobinPolicy("DC1"))  
    ).build();
```

Policy that decides how often the reconnection to a dead node is attempted.

```
Cluster cluster = Cluster.builder()
    .addContactPoints("10.158.02.40",
"10.158.02.44") .withRetryPolicy(DowngradingConsistencyRetryPolicy.INSTANCE)
    .withReconnectionPolicy(new ConstantReconnectionPolicy(1000))
    .withLoadBalancingPolicy(new TokenAwarePolicy(new DCAwareRoundRobinPolicy("DC1")))
    .build();
```

- ConstantReconnectionPolicy
- **ExponentialReconnectionPolicy (Default)**

Tracing is enabled on a per-query basis.

```
Query query = QueryBuilder
    .select()
    .all()
    .from("akeyspace", "user")
    .where(eq("username", "johnny"))
    .enableTracing();
```

```
ResultSet rs = session.execute(query);
ExecutionInfo executionInfo = rs.getExecutionInfo();
QueryTrace queryTrace = executionInfo.getQueryTrace();
```

Query Tracing

Connected to cluster: xerxes

Simplex keyspace and schema created.

Host (queried): /127.0.0.1

Host (tried): /127.0.0.1

Trace id: 96ac9400-a3a5-11e2-96a9-4db56cdc5fe7

activity	timestamp	source	source_elapsed
Parsing statement	12:17:16.736	/127.0.0.1	28
Pepraring statement	12:17:16.736	/127.0.0.1	199
Determining replicas for mutation	12:17:16.736	/127.0.0.1	348
Sending message to /127.0.0.3	12:17:16.736	/127.0.0.1	788
Sending message to /127.0.0.2	12:17:16.736	/127.0.0.1	805
Acquiring switchLock read lock	12:17:16.736	/127.0.0.1	828
Appending to commitlog	12:17:16.736	/127.0.0.1	848
Adding to songs memtable	12:17:16.736	/127.0.0.1	900
Message received from /127.0.0.1	12:17:16.737	/127.0.0.2	34
Message received from /127.0.0.1	12:17:16.737	/127.0.0.3	25
Acquiring switchLock read lock	12:17:16.737	/127.0.0.2	672
Acquiring switchLock read lock	12:17:16.737	/127.0.0.3	525
Appending to commitlog	12:17:16.737	/127.0.0.2	692
Appending to commitlog	12:17:16.737	/127.0.0.3	541
Adding to songs memtable	12:17:16.737	/127.0.0.2	741
Adding to songs memtable	12:17:16.737	/127.0.0.3	583
Enqueuing response to /127.0.0.1	12:17:16.737	/127.0.0.3	751
Enqueuing response to /127.0.0.1	12:17:16.738	/127.0.0.2	950
Message received from /127.0.0.3	12:17:16.738	/127.0.0.1	178
Sending message to /127.0.0.1	12:17:16.738	/127.0.0.2	1189
Message received from /127.0.0.2	12:17:16.738	/127.0.0.1	249
Processing response from /127.0.0.3	12:17:16.738	/127.0.0.1	345
Processing response from /127.0.0.2	12:17:16.738	/127.0.0.1	377

User Defined Types - Retrieval

```
USE ks;

CREATE TYPE address (
    street text,
    city text,
    zip int
);

CREATE TABLE user_profiles (
    email text PRIMARY KEY,
    address address
);
```

```
Row row = session.execute(
    "SELECT * FROM user_profiles").one();
UDTValue address =
row.getUDTValue("address");
```

```
// You get a map-like object with the
usual
// getters and setters (by name, by
index):
```

```
String street =
address.getString("street");
int zip       = address.getInt(2);
```

User Defined Types - Creation

```
USE ks;

CREATE TYPE address (
    street text,
    city text,
    zip int
);

CREATE TABLE user_profiles (
    email text PRIMARY KEY,
    address address
);
```

// Must first get hold of the datatype:
 // From an existing value:
 UserType addressType = address.getType();

// OR...

// From the cluster metadata:
 UserType addressType =
 cluster.getMetadata()
 .getKeyspace("ks")
 .getUserType("address");

User Defined Types - Creation

```
USE ks;

CREATE TYPE address (
    street text,
    city text,
    zip int
);

CREATE TABLE user_profiles (
    email text PRIMARY KEY,
    address address
);
```

```
UDTValue address =
    addressType.newValue()
        .setString("street", "...")
        .setString("city",
            "Washington")
        .setInt("zip", 20500);

// Now use it in a query like any other
type:
session.execute(
    "INSERT INTO user_profiles (email,
address) VALUES (?, ?)",
    "xyz@example.com", address);
```

Tuples - retrieval

```
USE ks;  
  
CREATE TABLE  
points_of_interest (  
    id int PRIMARY KEY,  
    name text,  
    coordinates  
        tuple<float, float>  
);
```

```
Row row = session.execute(  
    "SELECT * FROM  
points_of_interest").one();  
TupleValue coordinates =  
    row.getTupleValue("coordinates");
```

```
float latitude  = coordinates.getFloat(0);  
float longitude = coordinates.getFloat(1);
```


Tuples - creation

```
USE ks;

CREATE TABLE
points_of_interest (
  id int PRIMARY KEY,
  name text,
  coordinates
    tuple<float, float>
);
```

```
TupleType coordinatesType =
  TupleType.of(DataType.cfloat(),
    DataType.cfloat());

TupleValue newCoordinates =
  coordinatesType.newValue(48.858222F,
    2.2945F);
```

Query results ↔ Java objects

- Avoid boilerplate for common use cases (CRUD, etc.)
- Keep it **simple** and close to the bare metal
 - **Do not hide Cassandra from the developer**
 - Avoid "clever tricks" à la Hibernate

It comes as an additional Maven artifact:

```
<dependency>  
  <groupId>com.datastax.cassandra</groupId>  
  <artifactId>cassandra-driver-mapping</artifactId>  
  <version>2.1.0</version>  
</dependency>
```

Object Mapper Basic CRUD

```
USE ks;
```

```
CREATE TYPE address (  
    street text,  
    city text,  
    zip int  
);
```

```
CREATE TABLE user_profiles (  
    email text PRIMARY KEY,  
    address address  
);
```

```
@UDT(keyspace = "ks", name = "address")
```

```
public class Address {  
    private String street;  
    private String city;  
    private int zip;
```

```
// getters and setters omitted...
```

```
}
```

```
@Table(keyspace = "ks", name =  
"user_profiles")
```

```
public class UserProfile {  
    @PartitionKey  
    private String email;  
    private Address address;
```

```
// getters and setters omitted...
```

Company Confidential

Object Mapper Basic CRUD

```
USE ks;
```

```
CREATE TYPE address (  
    street text,  
    city text,  
    zip int  
);
```

```
CREATE TABLE user_profiles (  
    email text PRIMARY KEY,  
    address address  
);
```

```
MappingManager manager =  
    new MappingManager(session);
```

```
Mapper mapper =  
    manager.mapper(UserProfile.class);
```

```
UserProfile myProfile =  
    mapper.get("xyz@example.com");
```

```
ListenableFuture saveFuture =  
    mapper.saveAsync(anotherProfile);
```

```
mapper.delete("xyz@example.com");
```

Accessor custom queries

```
USE ks;

CREATE TYPE address (
    street text,
    city text,
    zip int
);

CREATE TABLE user_profiles (
    email text PRIMARY KEY,
    address address
);

@Accessor
interface ProfileAccessor {
    @Query("SELECT * FROM user_profiles
LIMIT :max")
    Result firstN(@Param("max") int limit);
}

ProfileAccessor accessor =
    manager.createAccessor(ProfileAccessor.class);

Result profiles = accessor.firstN(10);

// Result is like ResultSet, but specialized for a
// mapped class:
for (UserProfile profile : profiles) {
    System.out.println(
        profile.getAddress().getZip()
    );
}
```



Training Day | December 3rd

Beginner Track

- Introduction to Cassandra
- Introduction to Spark, Shark, Scala and Cassandra

Advanced Track

- Data Modeling
- Performance Tuning

Conference Day | December 4th

Cassandra Summit Europe 2014 will be the single largest gathering of Cassandra users in Europe. Learn how the world's most successful companies are transforming their businesses and growing faster than ever using Apache Cassandra.

<http://bit.ly/cassandrasummit2014>

EUROPE'S LARGEST GATHERING OF CASSANDRA DEVELOPERS

DECEMBER 3 - 4, 2014 | LONDON, U.K. | THE PARK PLAZA RIVERBANK HOTEL

Twitter: [@CassandraEurope](#) | [#CassandraSummit](#)